

computer systems performance

lecture 2

memory hierarchy & CPU parallelism

Pınar Tözün

February 4, 2026

borrowed some slides / inspiration from

A. Ailamaki, E. Liarou, P. Tözün, D. Porobic, I. Psaroudakis

[How to Stop Under-Utilization and Love Multicores.](#)

ICDE 2015 Tutorial

agenda

lecture

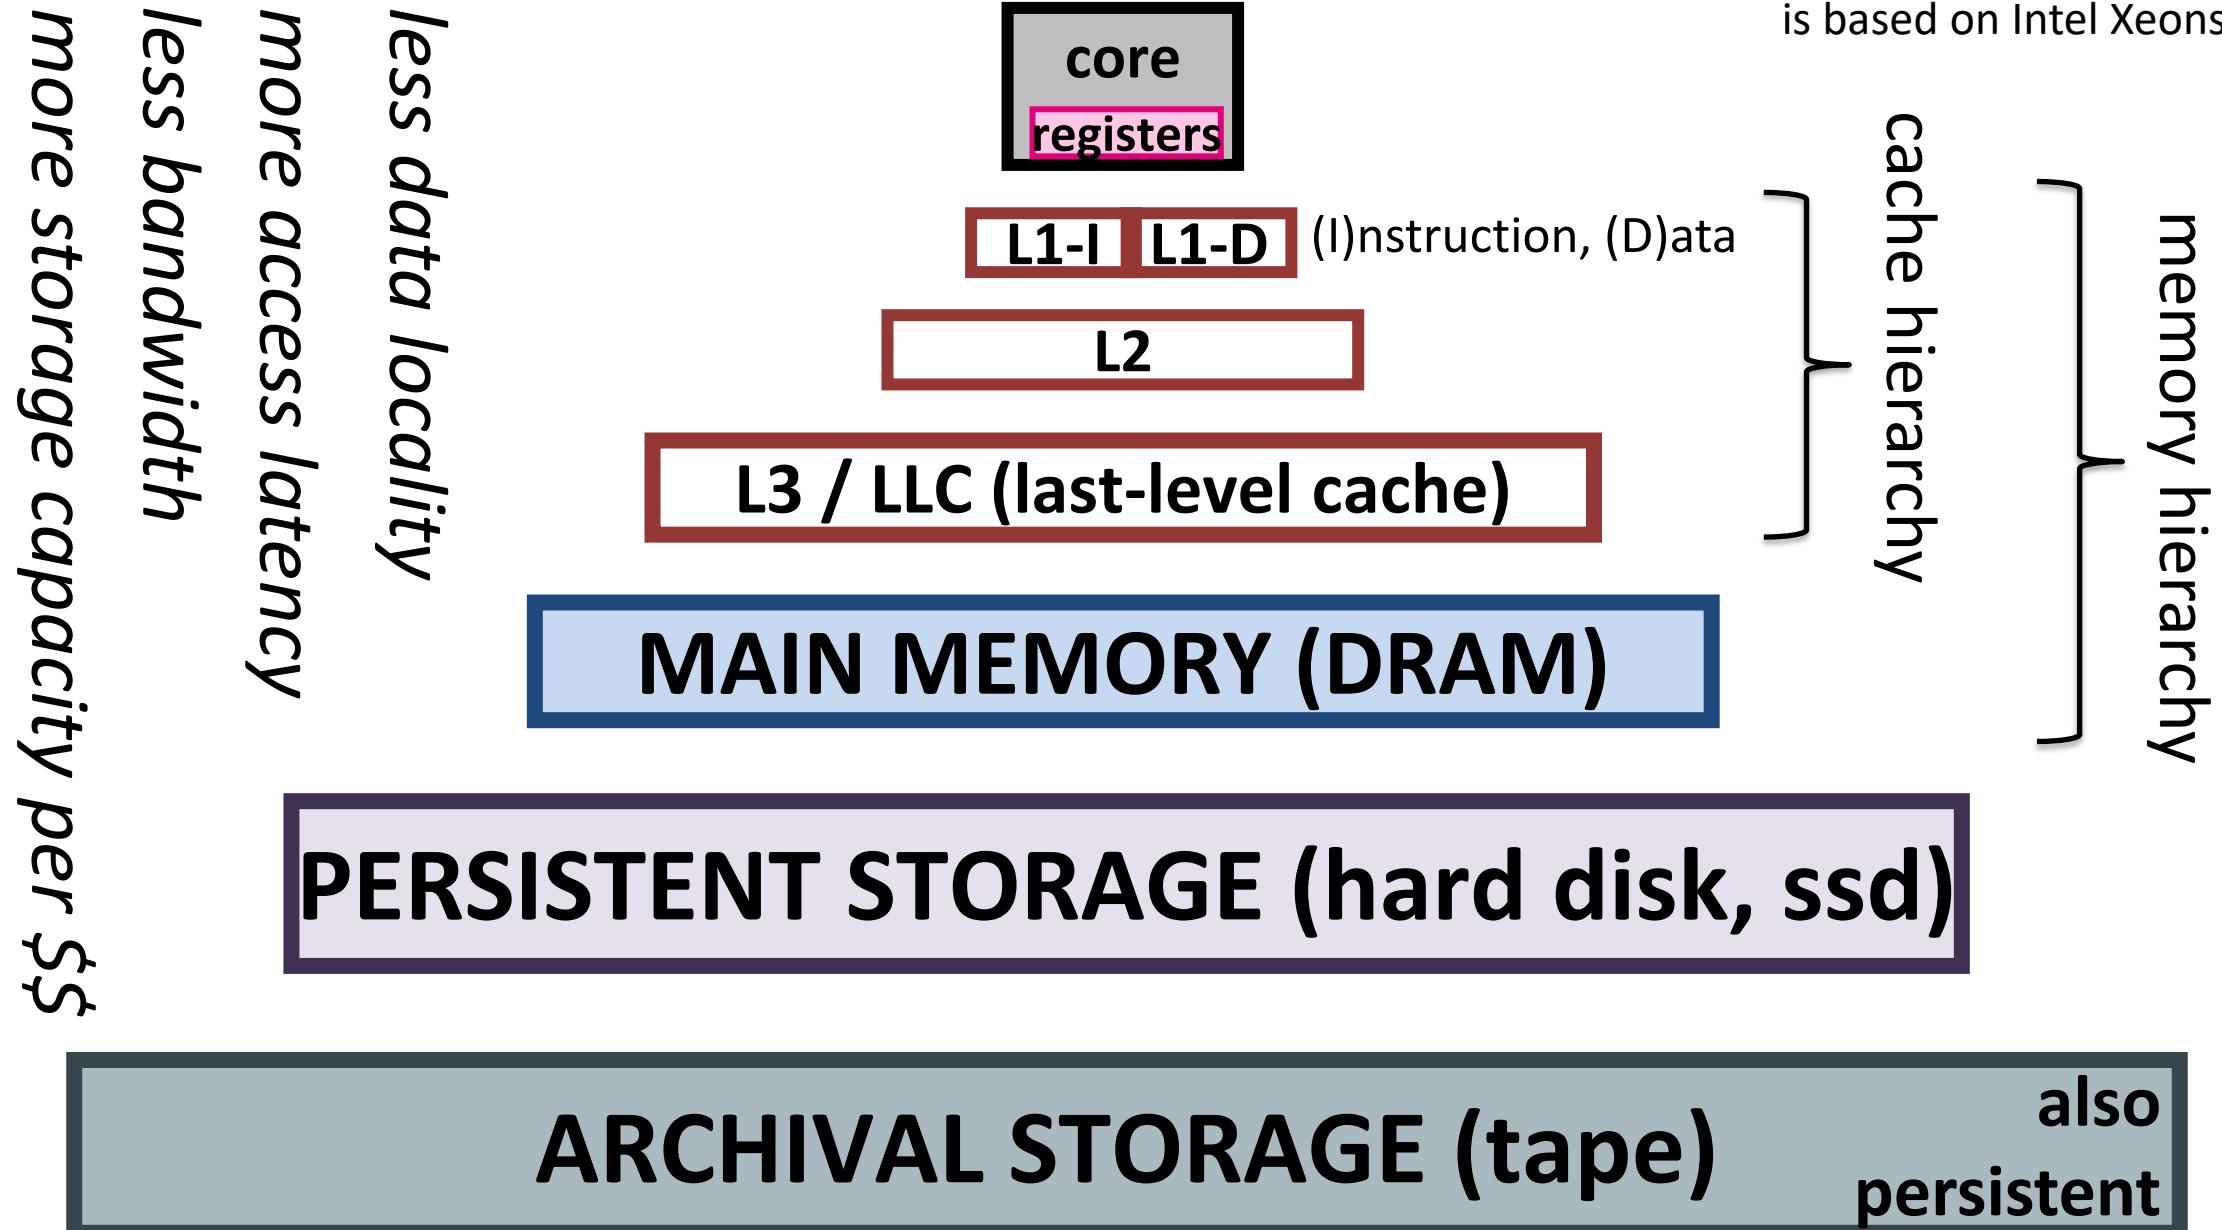
- storage hierarchy overview
- local memory hierarchy: DRAM & caches
- hardware parallelism: implicit & explicit
- different memory hierarchies

exercises

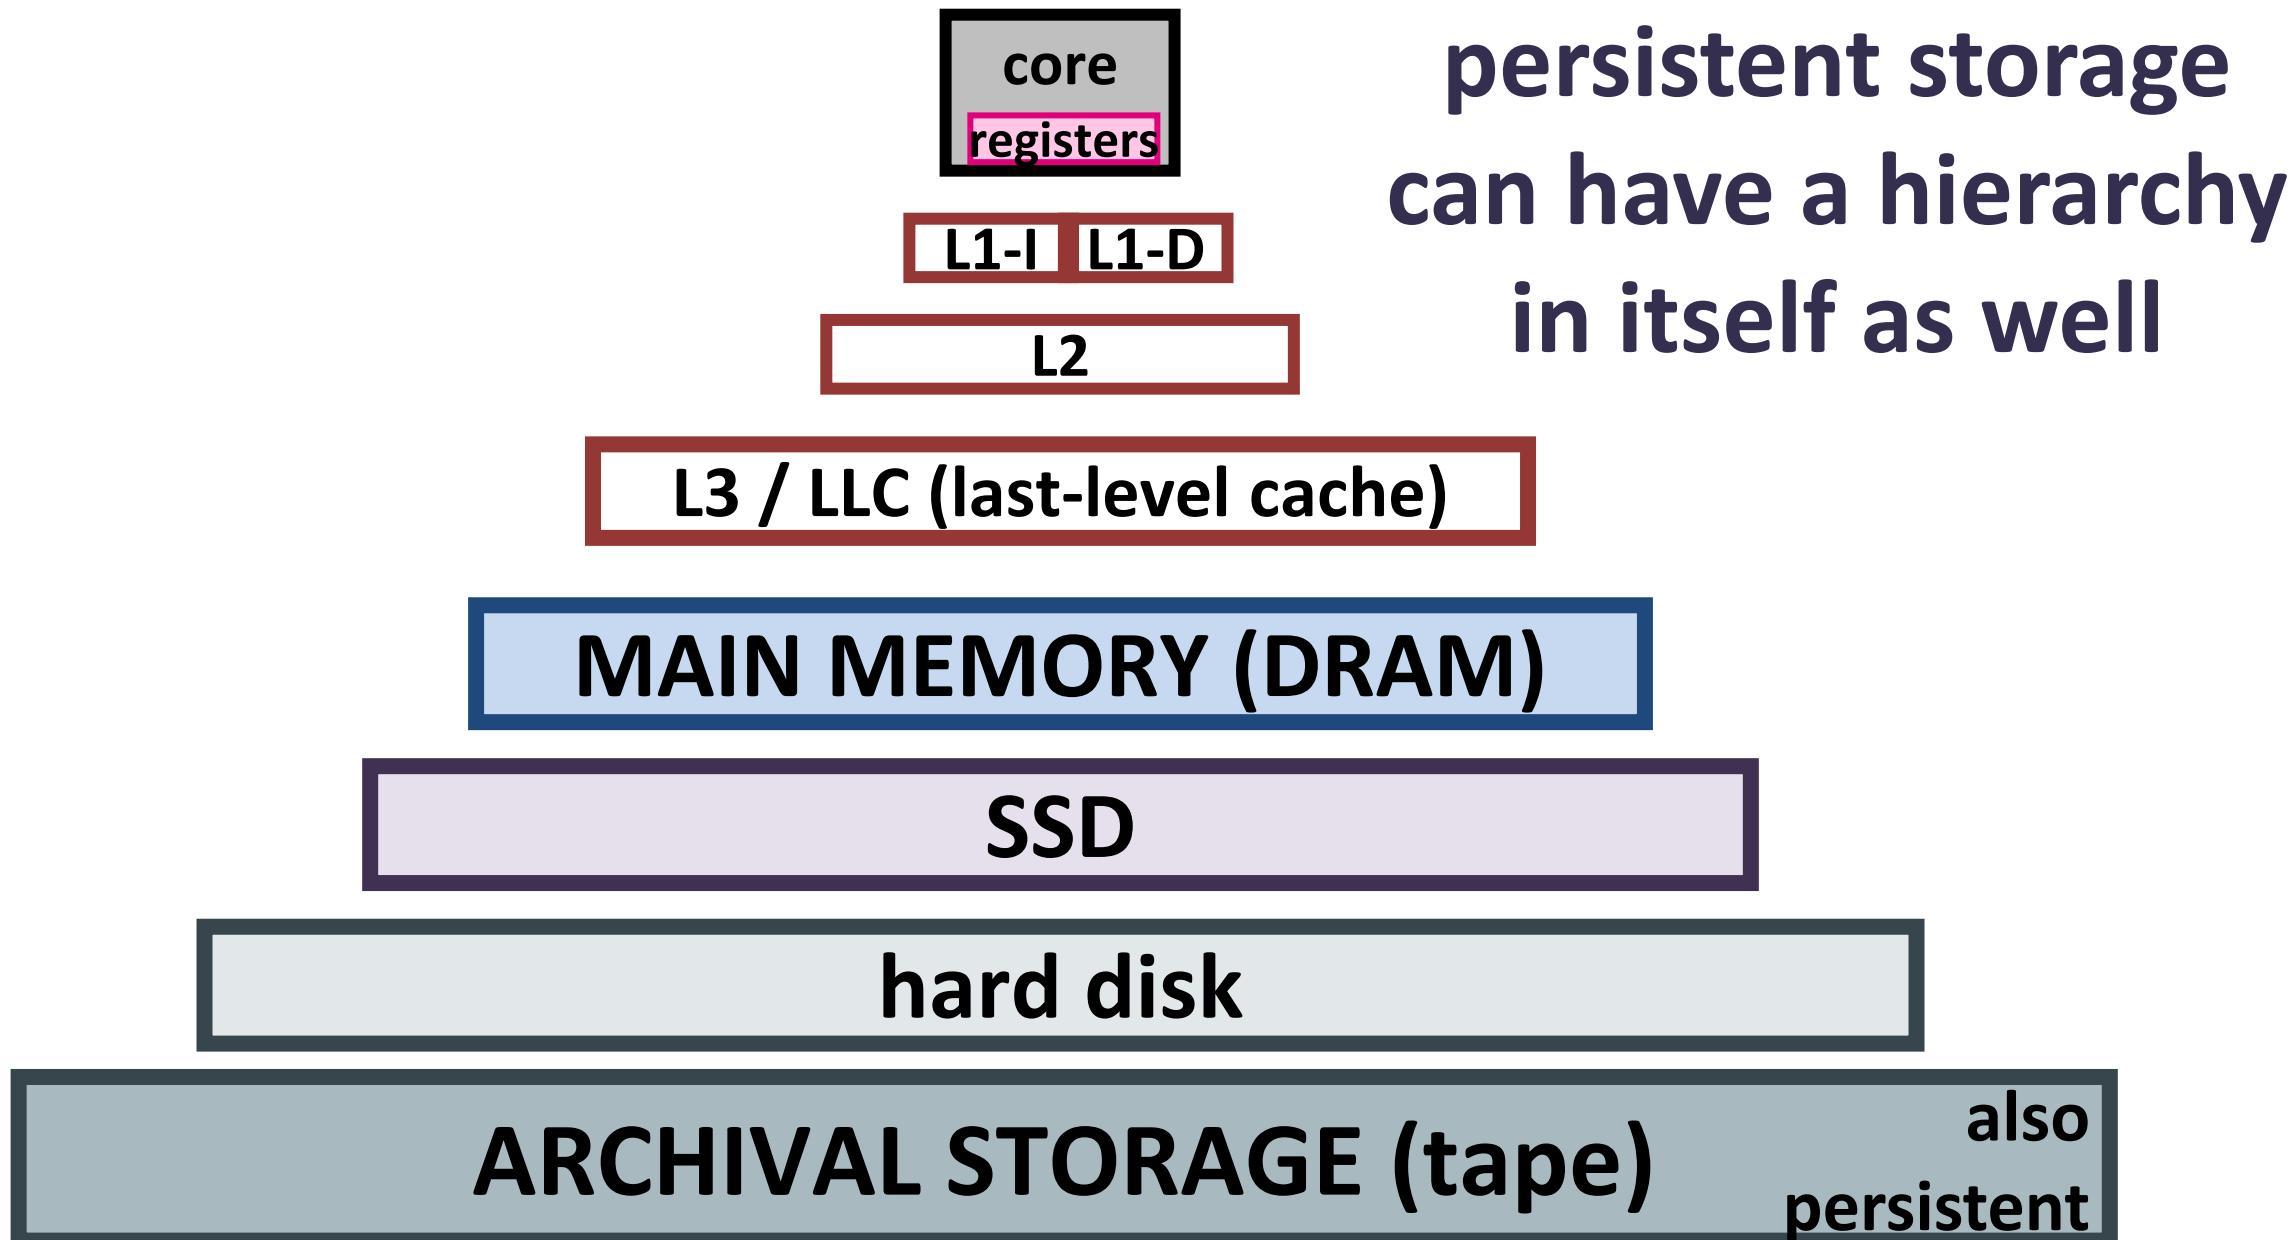
- paper discussion
- c/cpp examples
- setting core affinity

(typical) storage hierarchy

disclaimer: memory hierarchy
is based on Intel Xeons' here



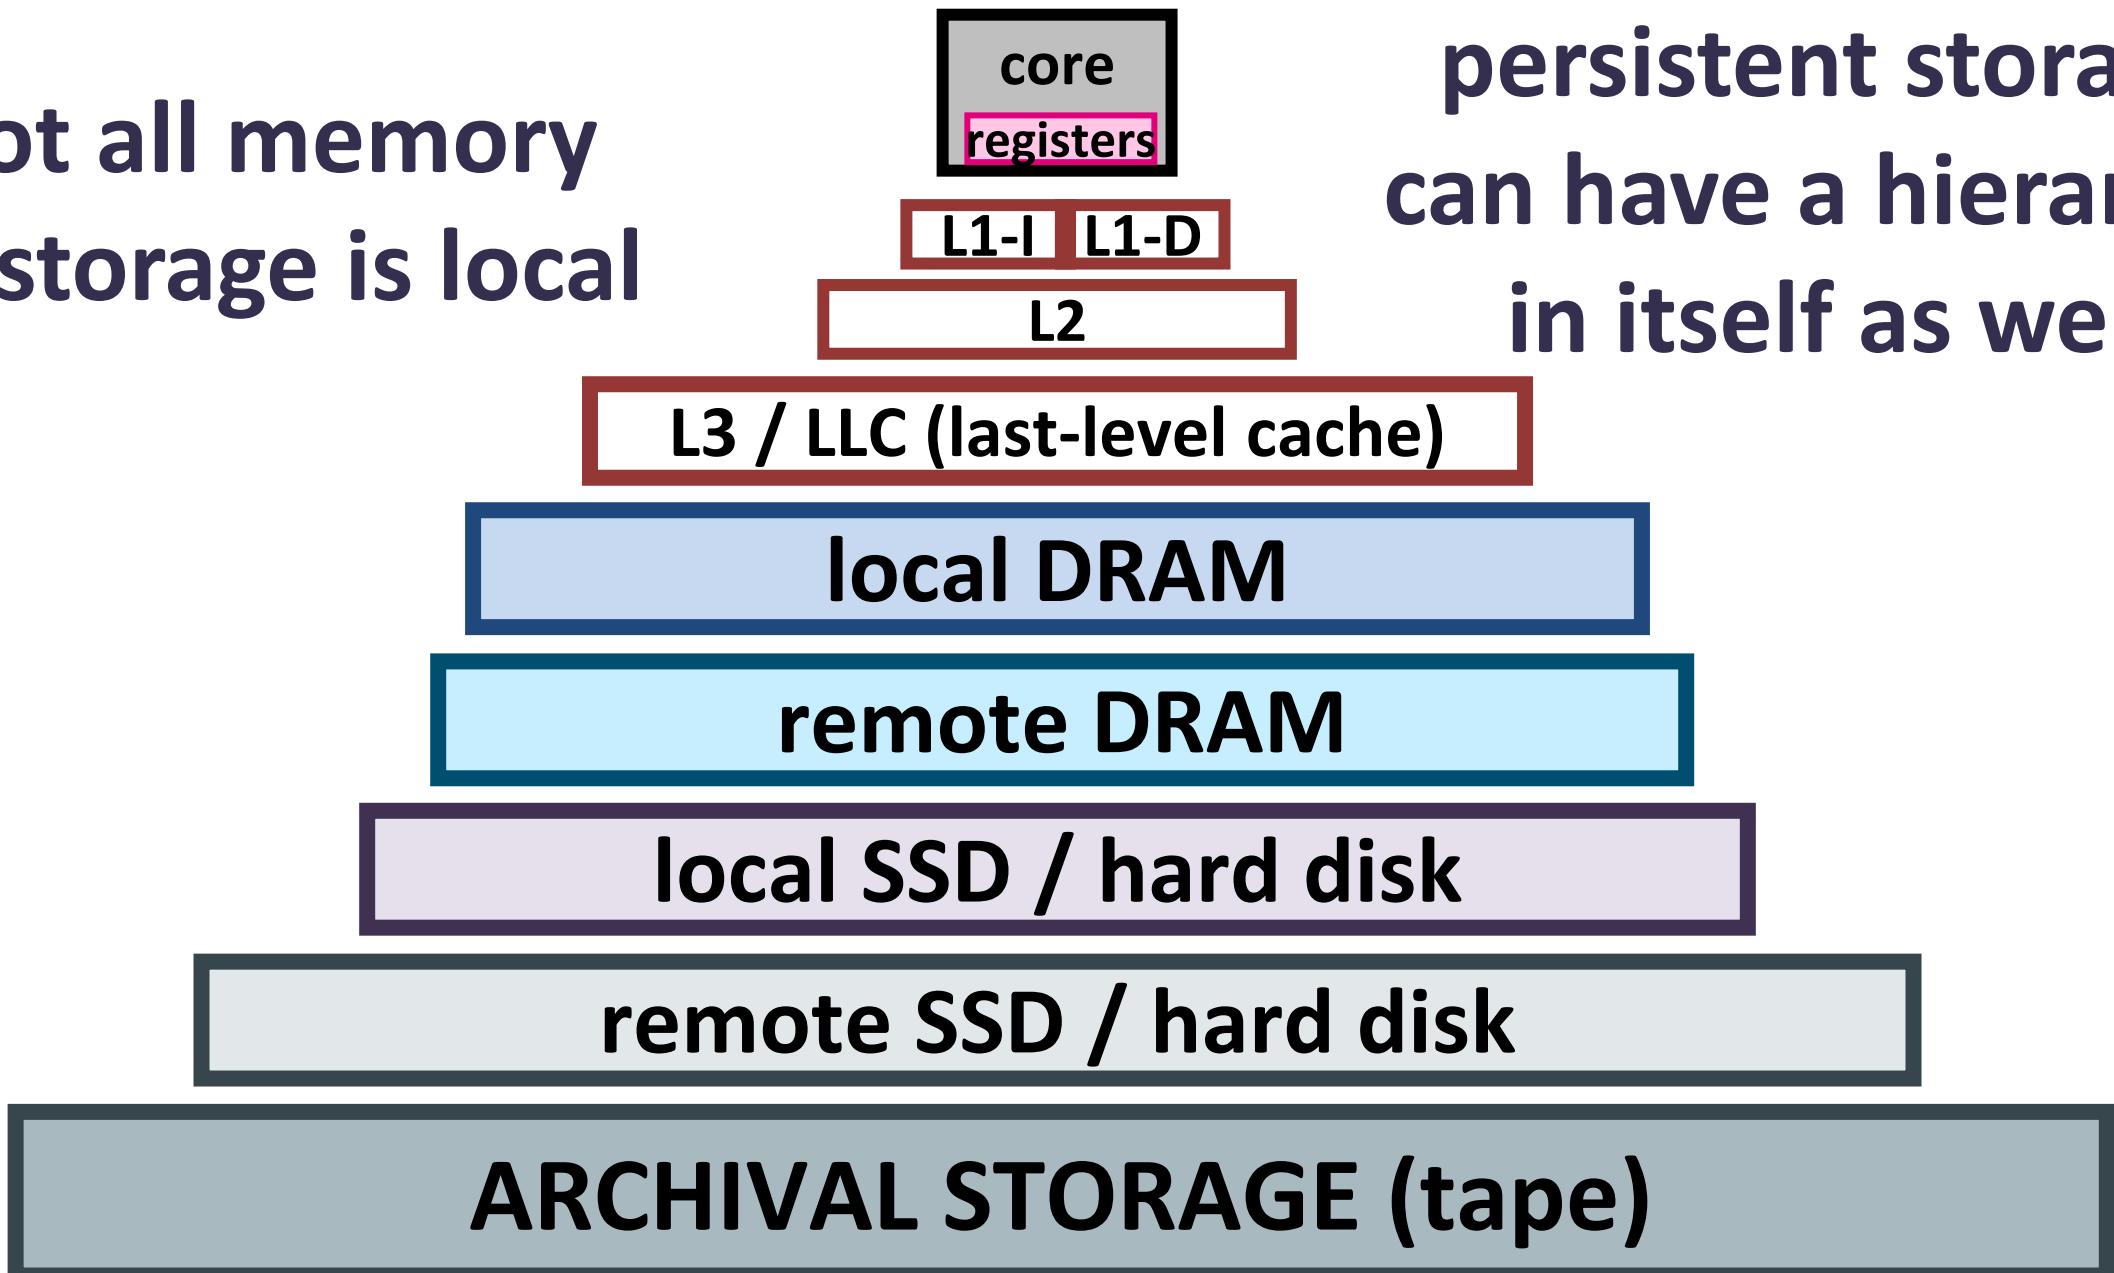
(typical) storage hierarchy



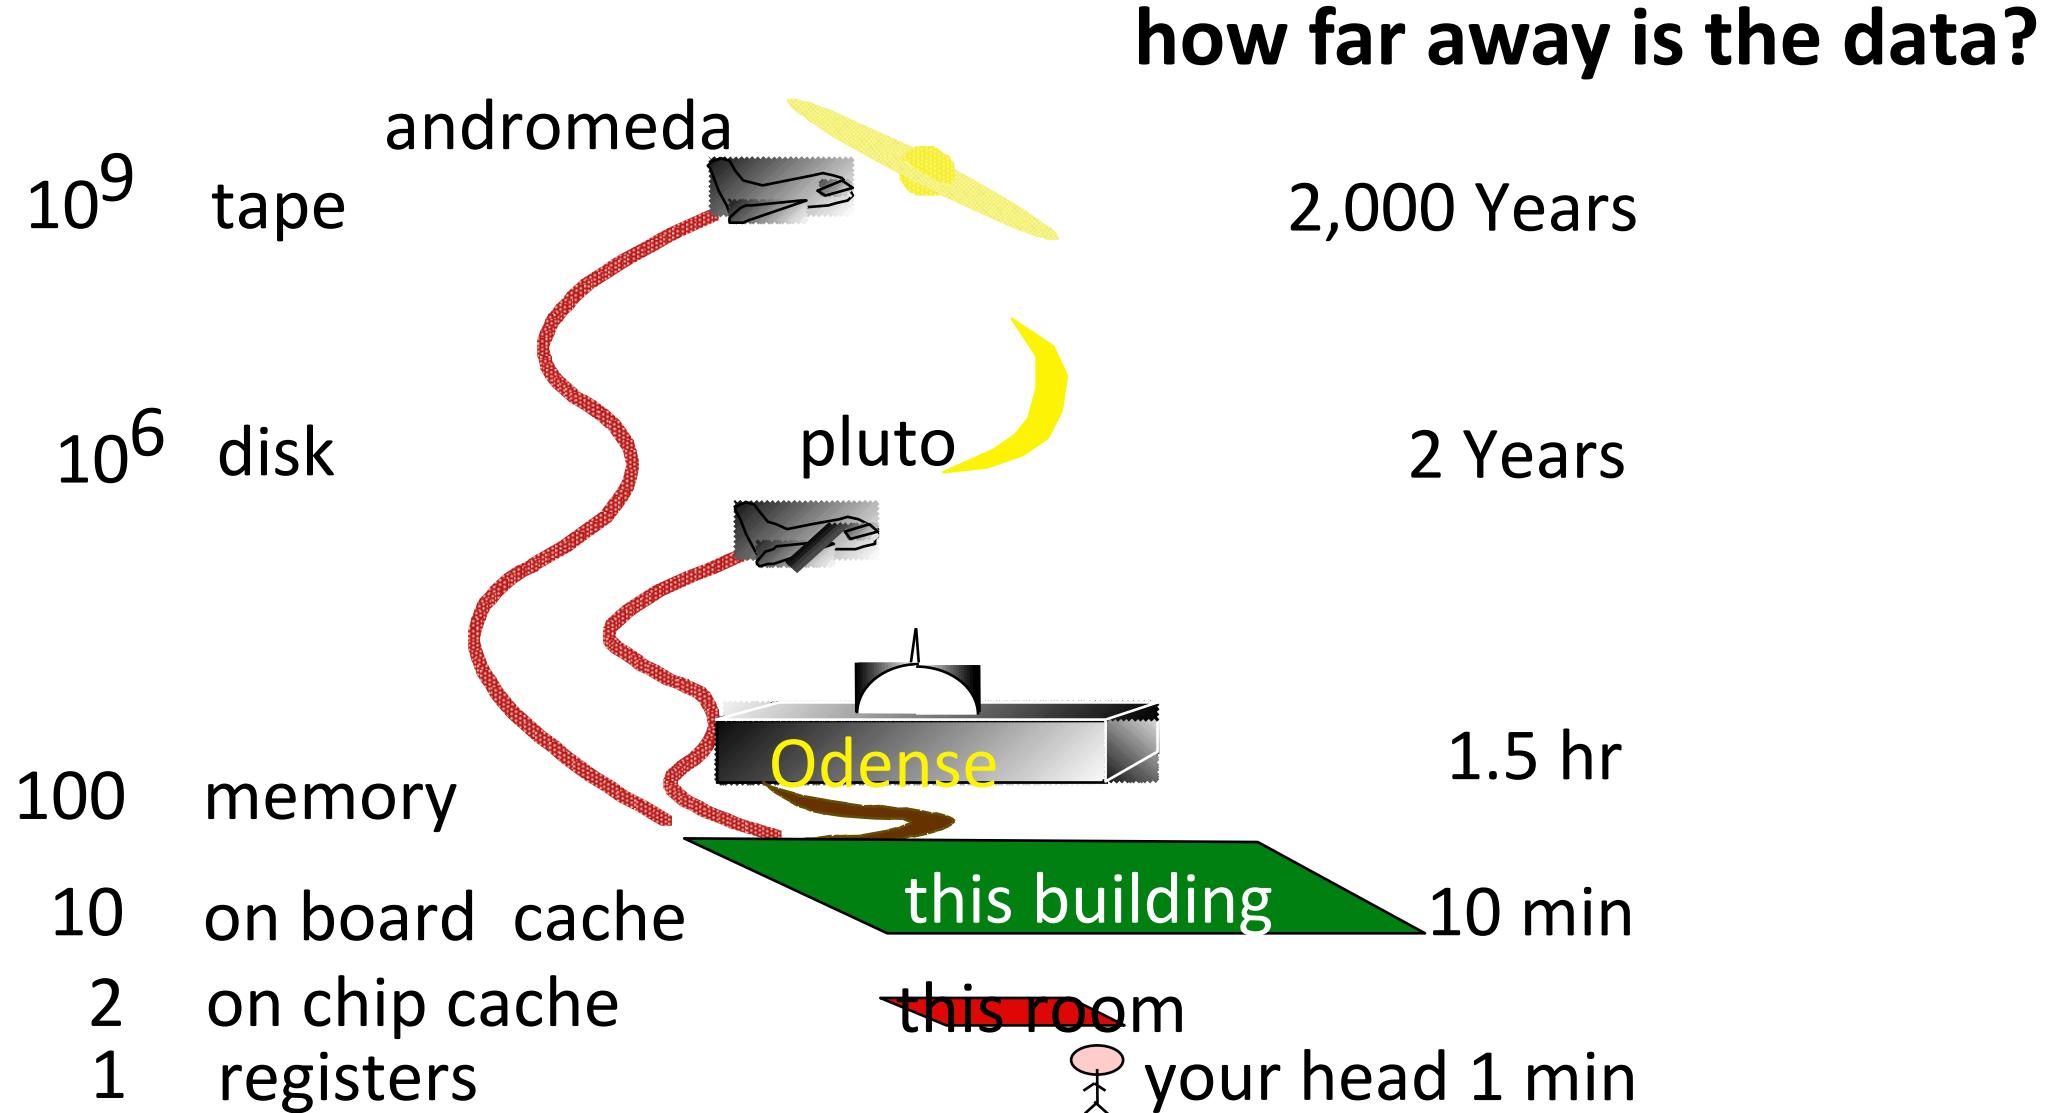
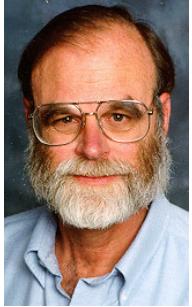
(typical) storage hierarchy

not all memory
or storage is local

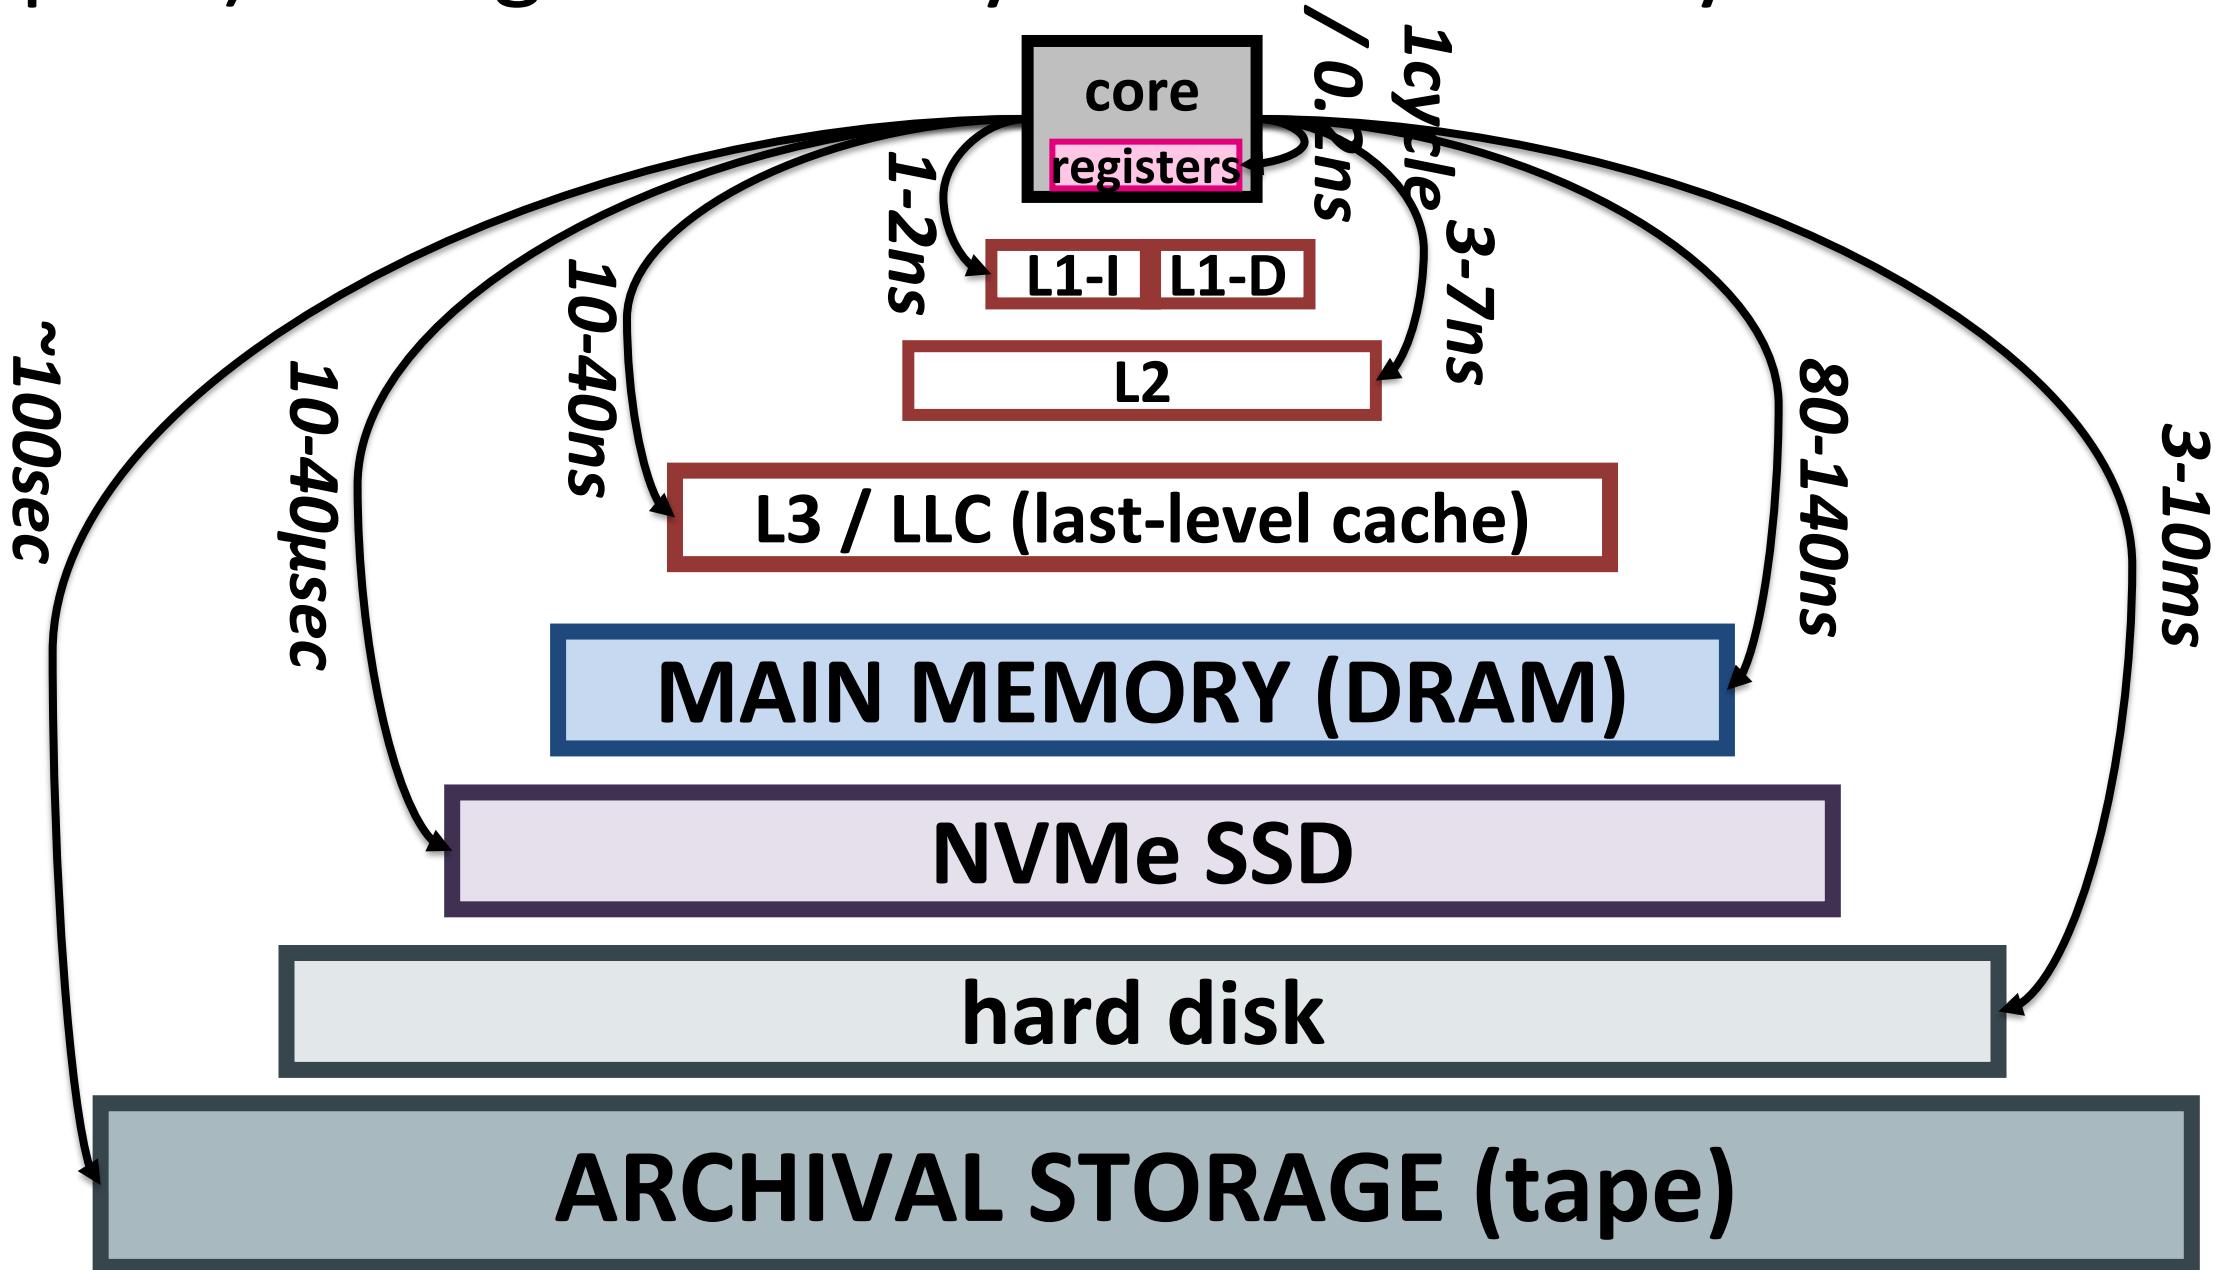
persistent storage
can have a hierarchy
in itself as well



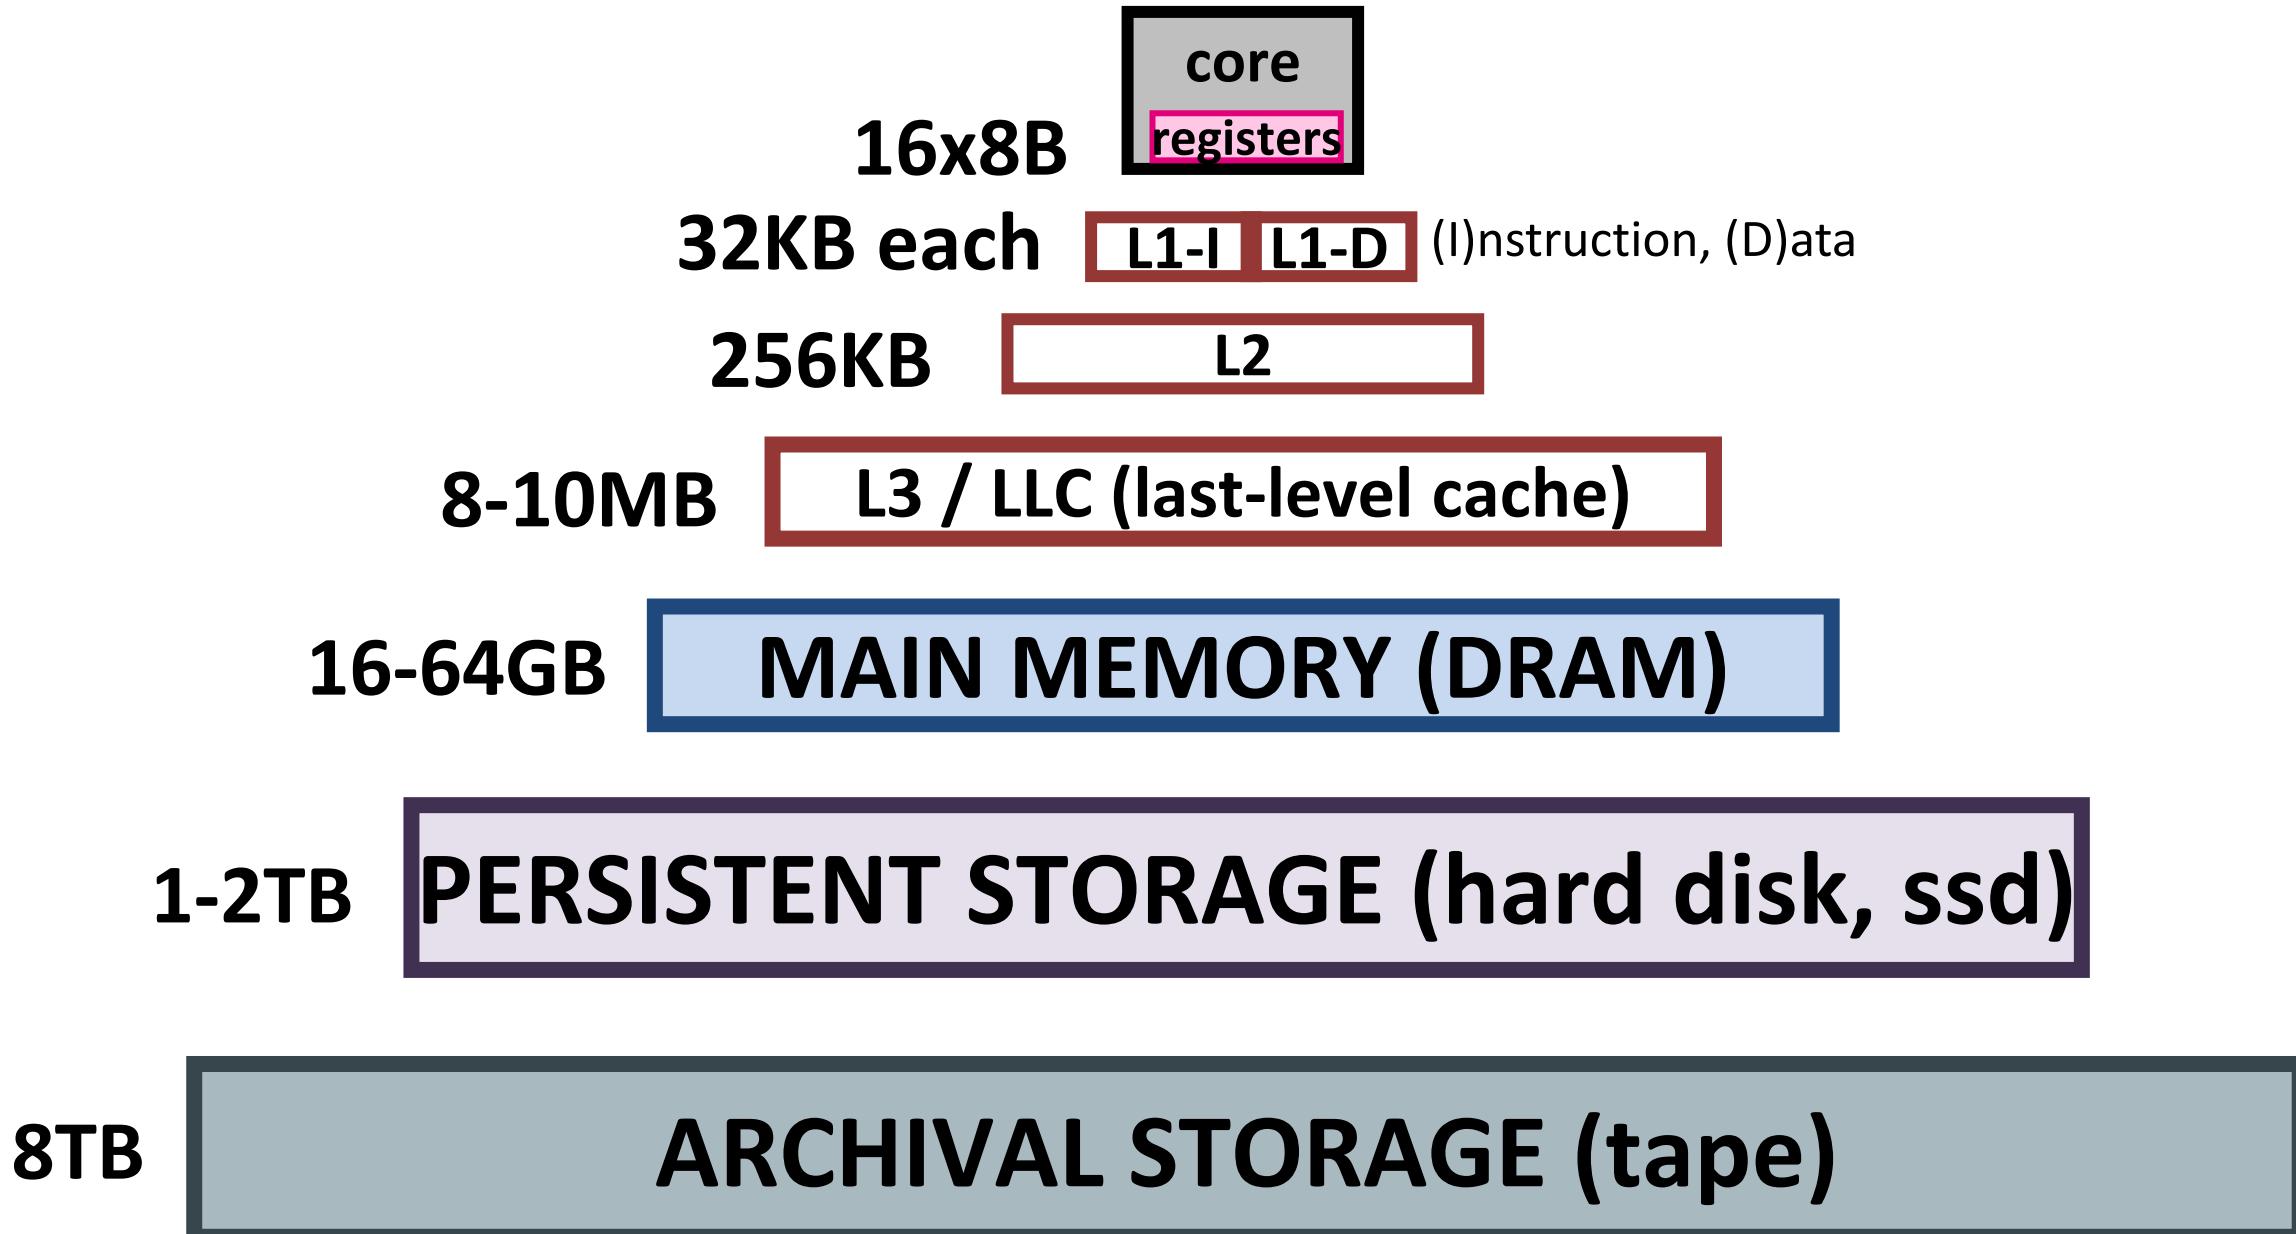
Jim Gray's storage latency analogy



(typical) storage hierarchy – access latency

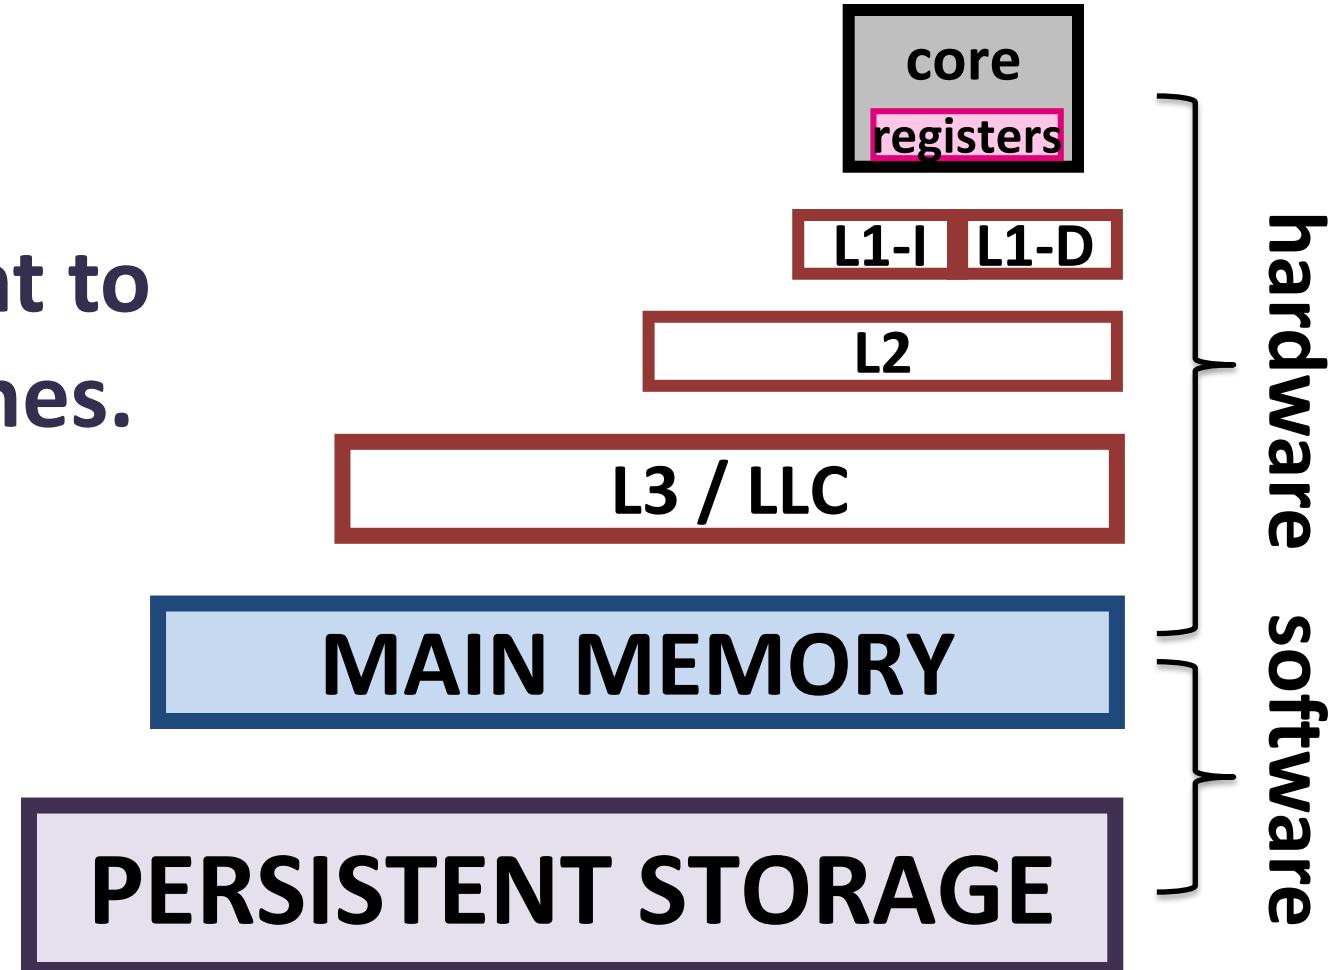


(typical) storage hierarchy – capacity



(typical) storage hierarchy – capacity

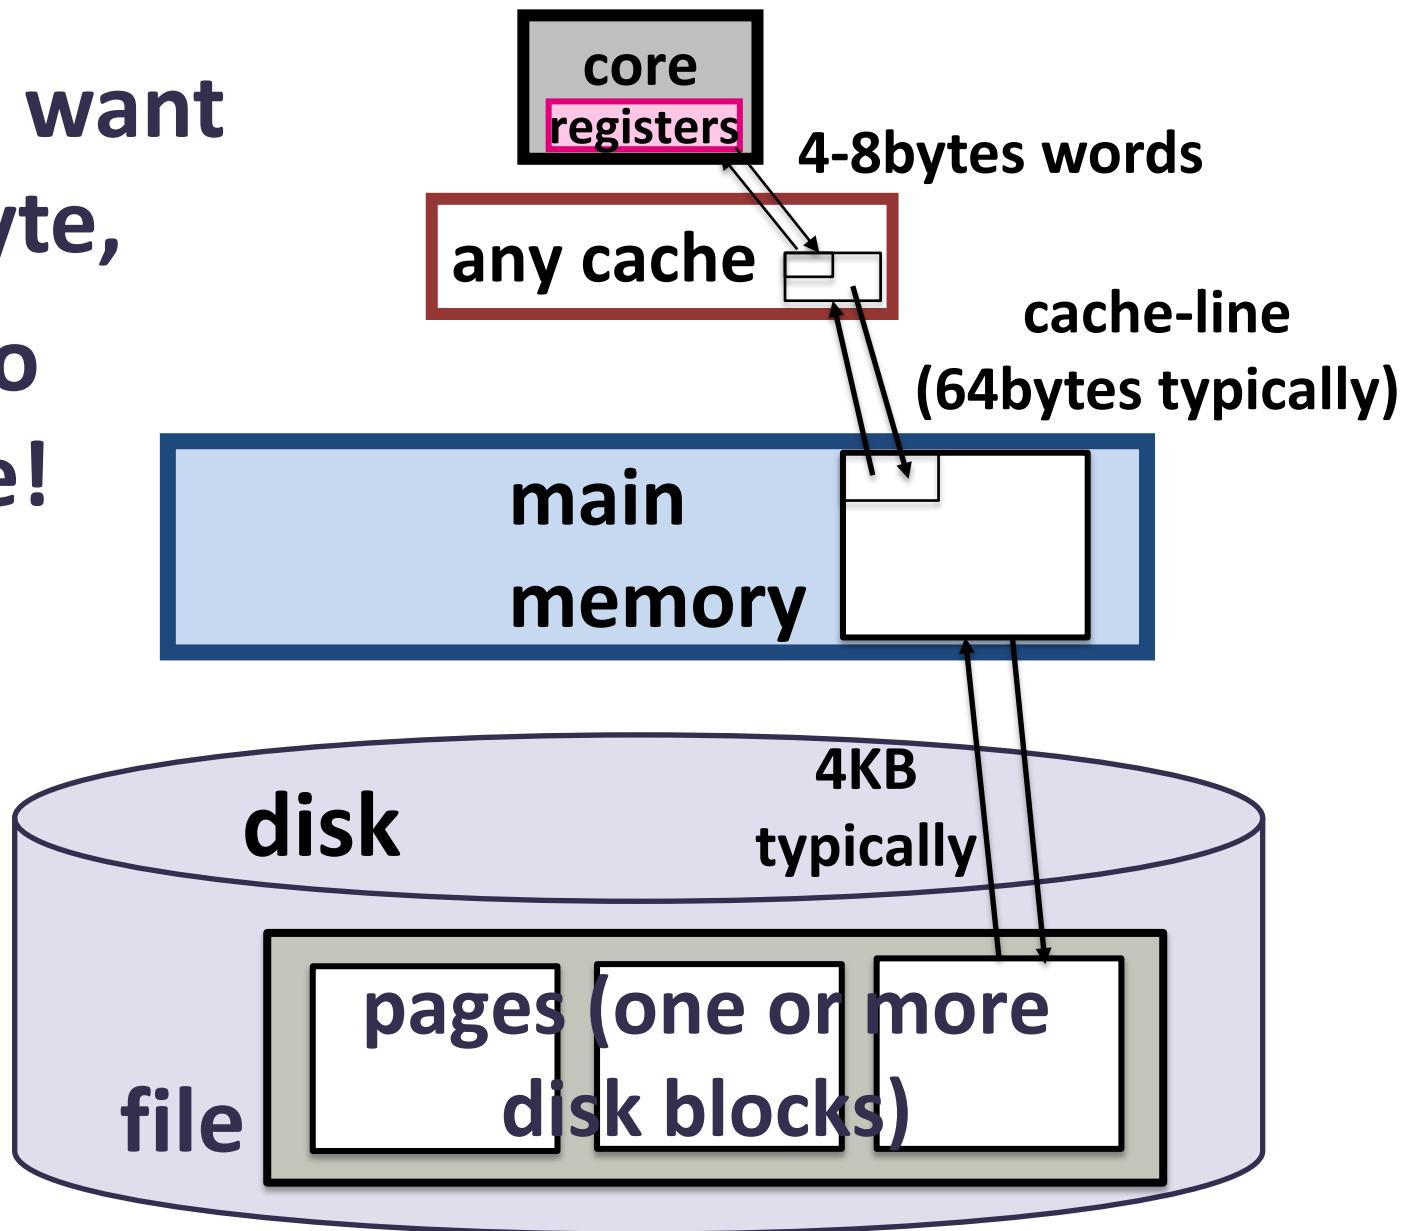
means that we do not write code to explicitly manage data movement to registers, L1, L2, L3 caches.



but we do this for moving data from/to persistent storage.

movement of data in storage hierarchy

even if you want
to read 1byte,
you need to
move more!

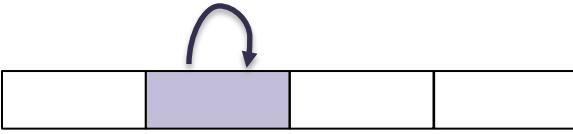


principle of locality

programs tend to use data & instructions with addresses near or equal to those they have used recently

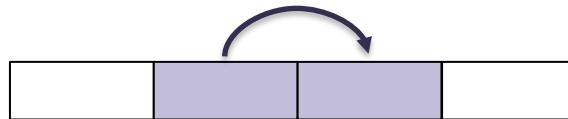
temporal locality

recently referenced items are likely to be referenced again in near future



spatial locality

items with nearby addresses tend to be referenced close together in time



impacts design decisions for both software & hardware

principle of locality – example

```
for (int i = 0; i < array.size(); i++) {  
    sum += array[i];  
}
```

data

- **temporal** → referencing *sum* in each iteration
- **spatial** → iterating over array element

instructions

- **temporal** → looping over the same code
- **spatial** → executing instructions in sequence

caching

goal is to increase locality for cores

to reduce access latency for frequently accessed data

higher levels cache data from lower levels

inclusivity

- most hardware vendors build inclusive cache hierarchy
- software controlled caching can be more complex
 - e.g., some main-memory optimized database systems don't persist indexes

data replacement when no space left

replacement policy can be a crucial optimization

dealing with modified data

write through – immediately or

write back – at replacement time

agenda

lecture

- storage hierarchy overview
- local memory hierarchy: DRAM & caches
- hardware parallelism: implicit & explicit
- different memory hierarchies

exercises

- paper discussion
- c/cpp examples
- setting core affinity

random-access memory (RAM)

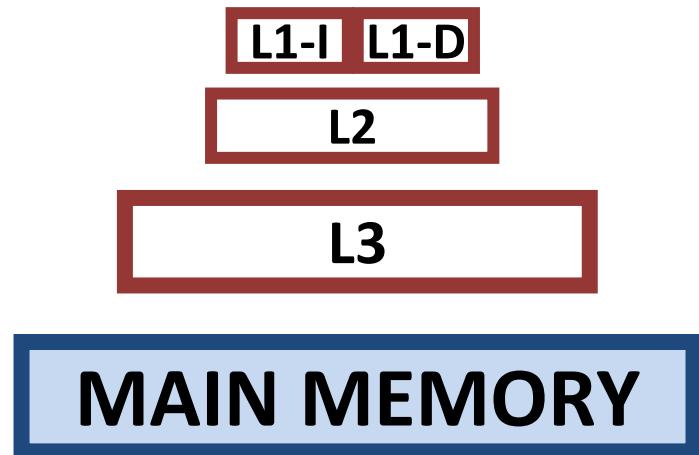
(almost) ***constant random-access latency*** wherever the data is

volatile → will loose data once power is lost

sequential access is slightly faster

because of ***prefetching***:

if you fetch a block from main-memory to caches,
hardware usually prefetches the adjacent block



D(dynamic)RAM vs S(static)RAM

- DRAM is most common for *main memory*, SRAM is used for *caches & registers*
- DRAM requires refresh, otherwise loses data even when power is on
- SRAM is more energy-efficient, but also more expensive,
so DRAM is preferred for main memory

caches

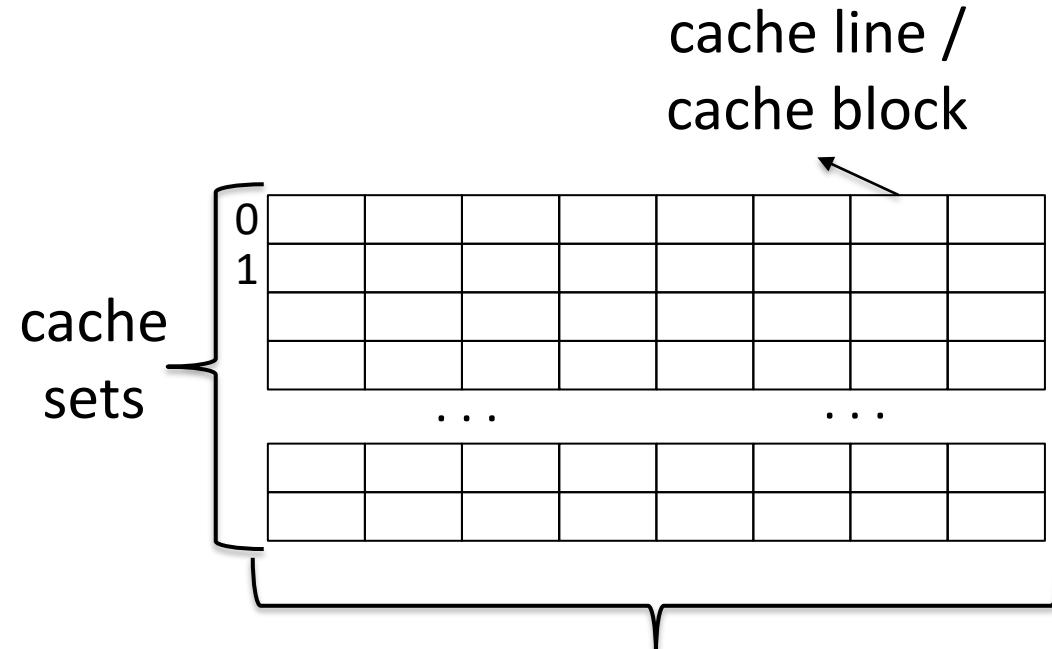
consider cache like a 2-dimentional array

example:

32KB 8-way L1 cache

64bytes cache lines

64bit memory addresses



$$32 * 1024 / 64 = 512 \text{ cache lines}$$

$$512 / 8 = 64 \text{ cache sets}$$

64bit address (a) decomposition for cache →

$$\text{TAG} = 52 \text{ bits} \mid \text{SET} = 6 \text{ bits} \mid \text{OFFSET} = 6 \text{ bits}$$

to determine if a
is cache hit/miss

which set to
look at

which byte to
read in cache line

cache misses

... when an item is not in the cache

compulsory misses

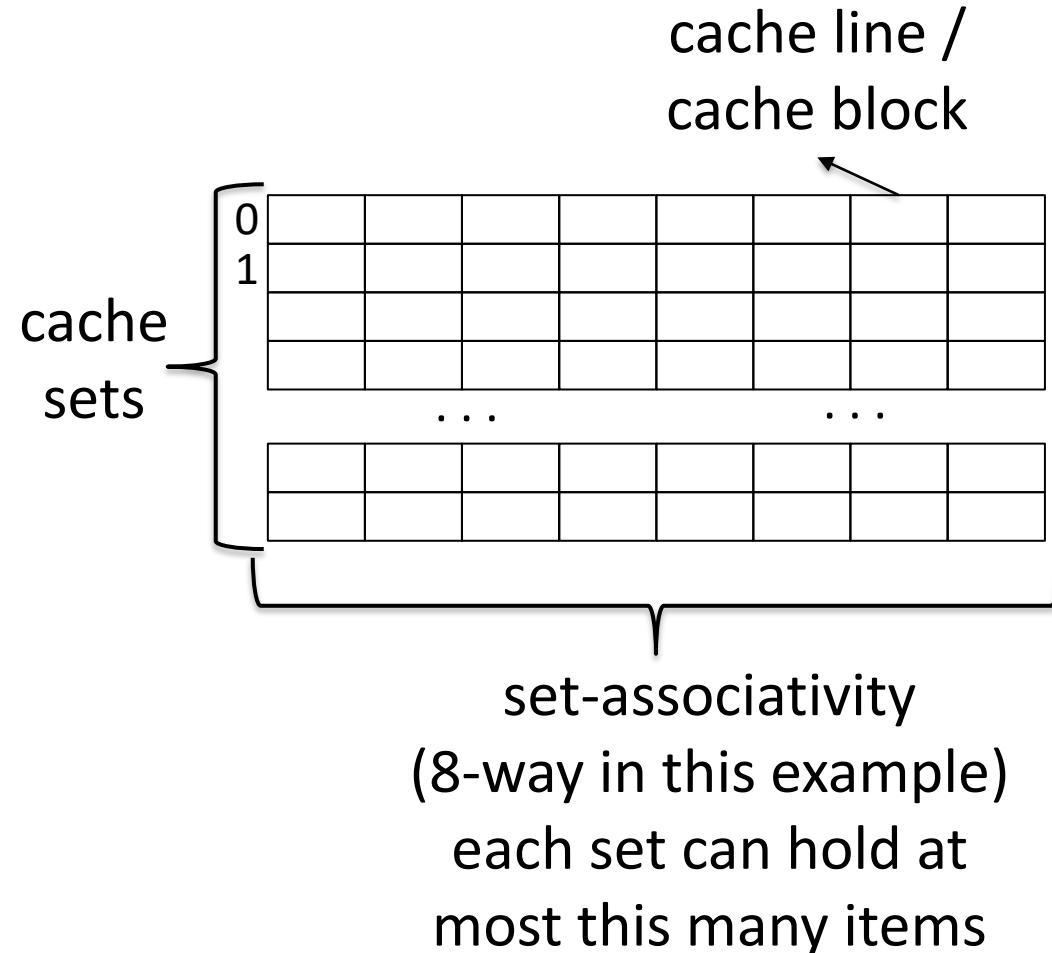
misses that occur because cache is empty

conflict misses

misses that occur due to a cache set being full even though cache has space
→ higher with lower associativity

capacity misses

misses due to cache being full
→ common when active data/instruction set size bigger than total cache size



hardware prefetching

a = address

- ***next-line*** prefetching: miss $a \rightarrow$ fetch $a+1$
- ***stream*** prefetching: miss $a, a+1 \rightarrow$ fetch $a+2, a+3$
- ✓ favors sequential access & spatial locality
- ✗ instructions: branches, function calls
 - branch prediction handles this to some extend
- ✗ data: pointer chasing – common for index accesses
 - ***stride*** prefetching handles this to some extend
 - miss $a, a+20 \rightarrow$ fetch $a+40, a+60$

**prefetchers on real hardware are kept simple
better accuracy requires more complexity**

summary: memory hierarchy

Different layers of the hierarchy have different characteristics & require different optimizations from the software side.

Principle of ***caching*** is to improve locality for frequently accessed data.

Sequential access is faster even in main-memory, where random and sequential access latency is the same, because of hardware prefetching.

It is important to be aware of the hierarchy to utilize it more effectively & have more efficient software layers.

goal: minimize data access latency!

agenda

lecture

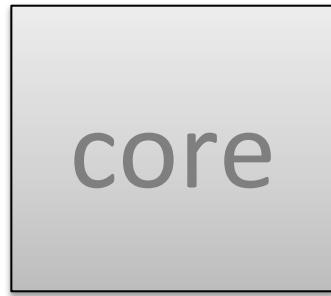
- storage hierarchy overview
- local memory hierarchy: DRAM & caches
- hardware parallelism: implicit & explicit
- different memory hierarchies

exercises

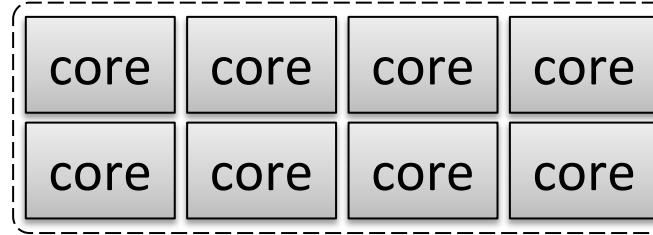
- paper discussion
- c/cpp examples
- setting core affinity

central processing unit (CPU) evolution

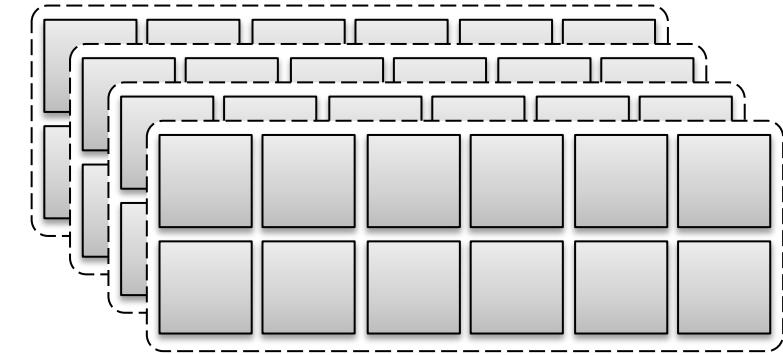
2005



single-core CPUs



multicore CPUs



multisocket
multicore CPUs

*faster & more-complex
cores over time*

*similar speed & complexity in a core,
more parallelism over time*

for Moore's law to be practical
you need Dennard scaling!

Moore's law

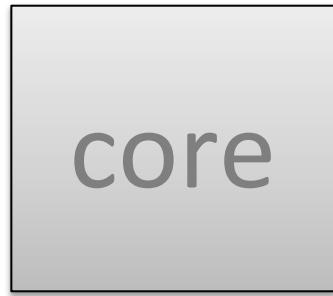
“... the observation that the number of transistors in a dense integrated circuit doubles approximately every two years.”

Dennard scaling

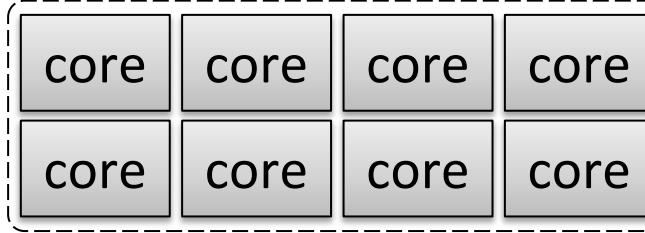
“ ... as transistors get smaller their power density stays constant, so that the power use stays in proportion with area: both voltage and current scale (downward) with length.”

central processing unit (CPU) evolution

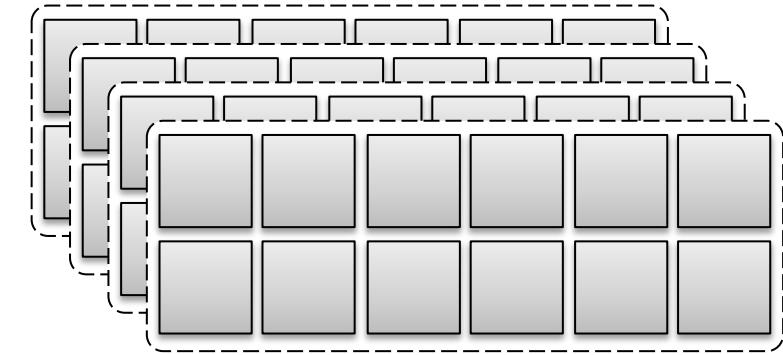
2005



single-core CPUs



multicore CPUs



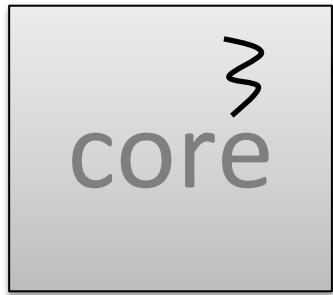
multisocket
multicore CPUs

around 2005, Dennard scaling stopped

switching to multicores kept Moore's Law alive

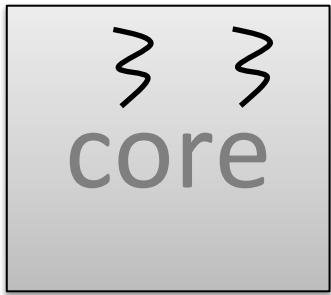
types of hardware parallelism

implicit/vertical parallelism



→ thread
what runs
your programs

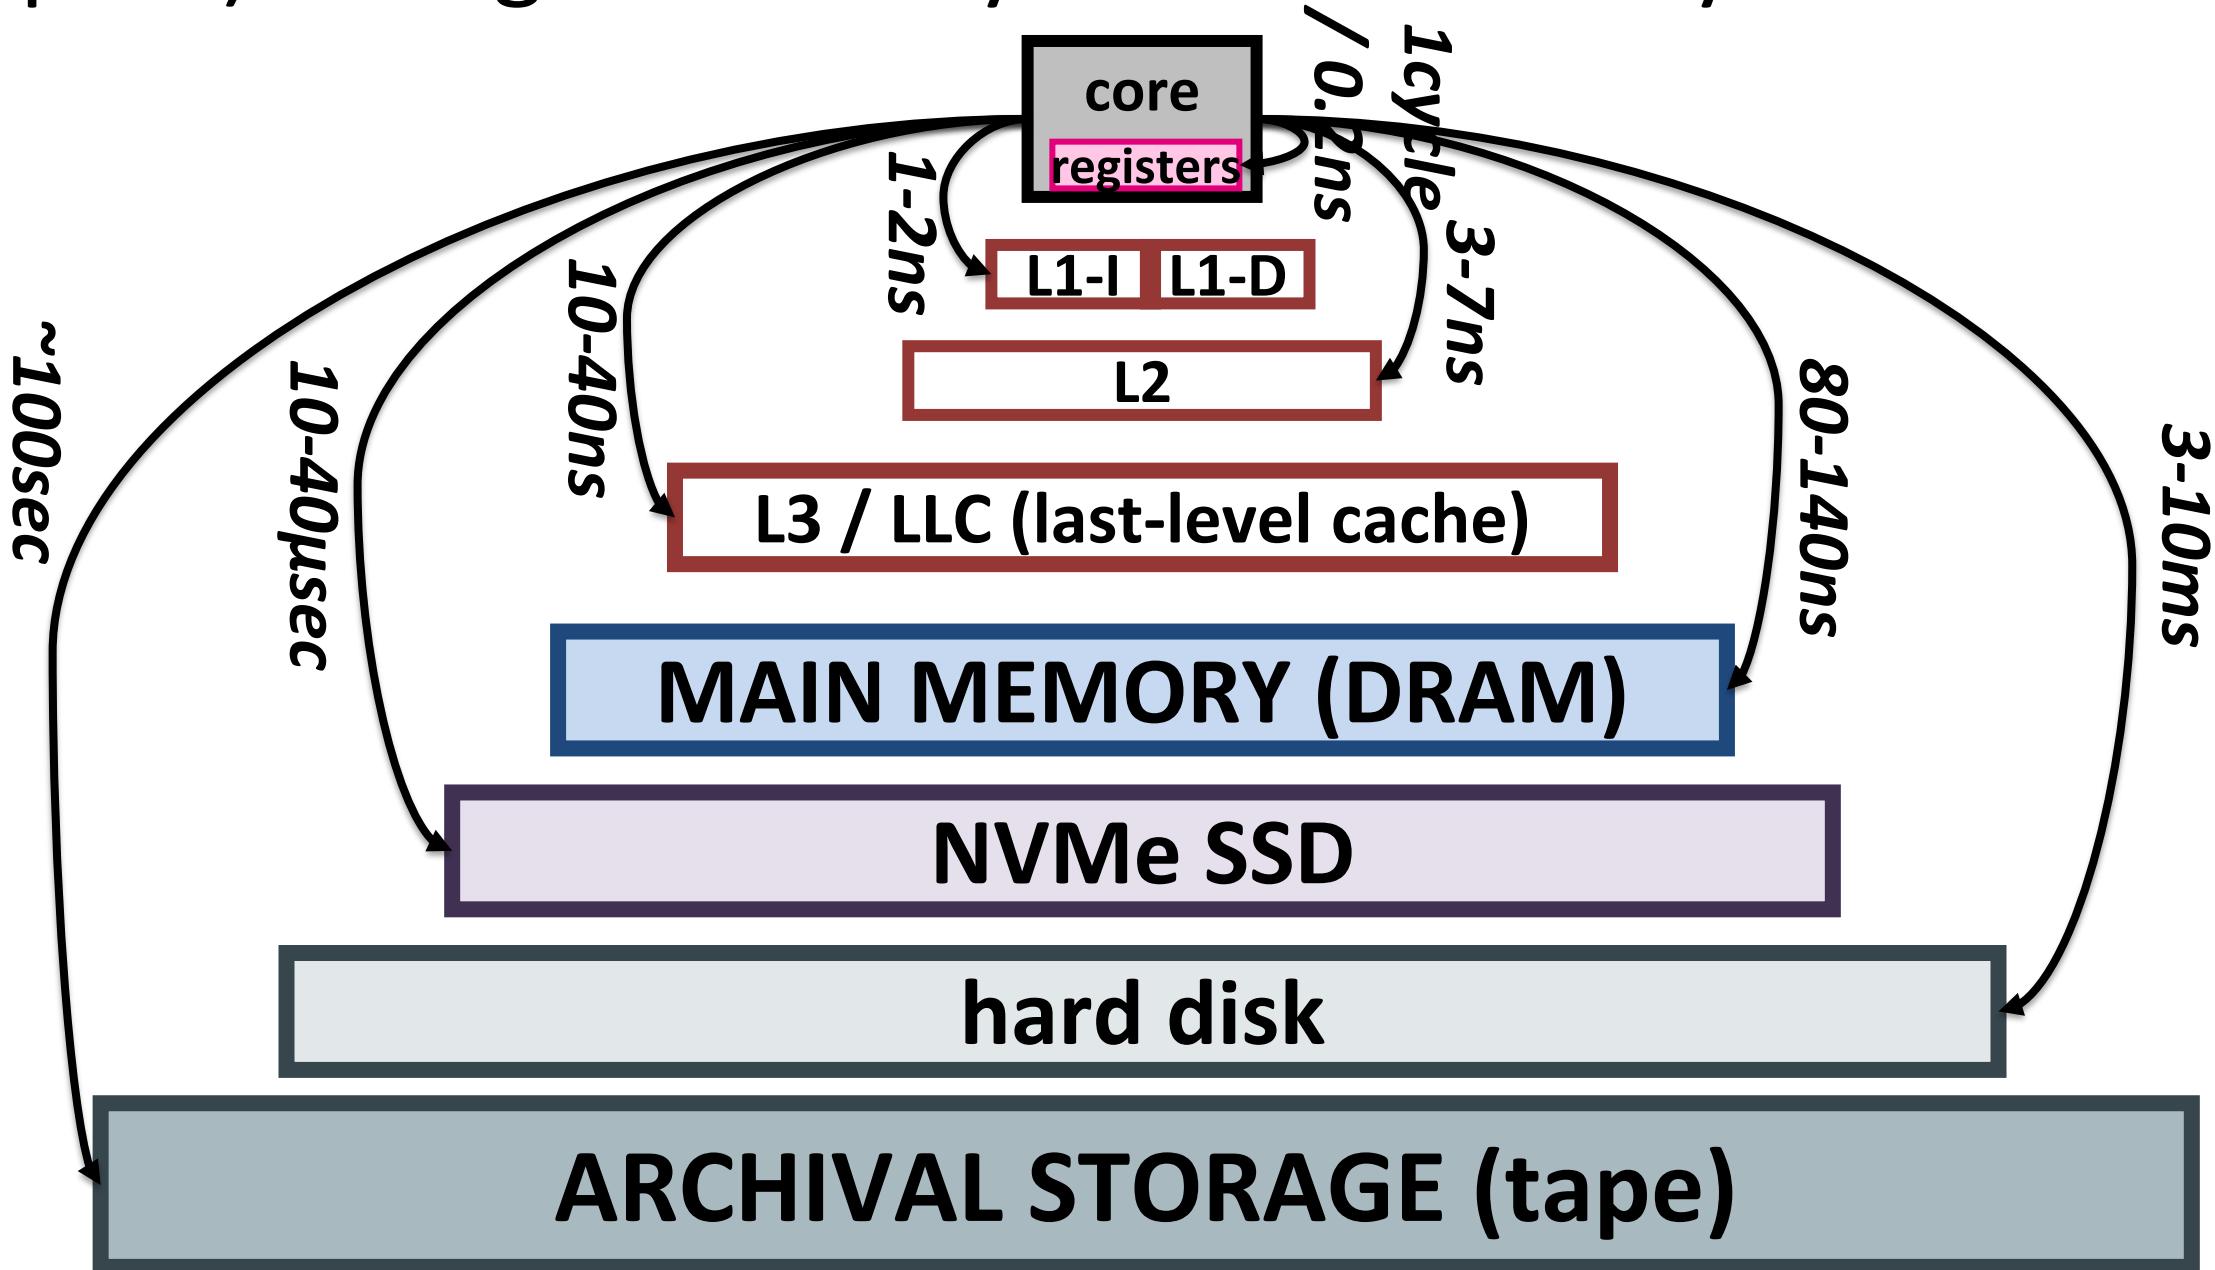
instruction & data parallelism
hardware does this automatically



multithreading
threads share
execution cycles
on the same core

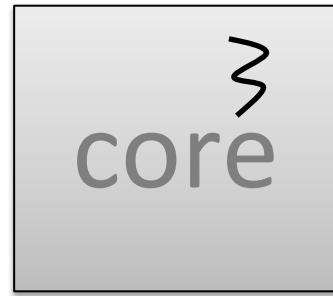
why do we need this?

(typical) storage hierarchy – access latency



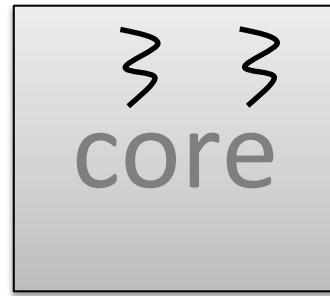
types of hardware parallelism

implicit/vertical parallelism



thread
what runs
your programs

instruction & data parallelism
hardware does this automatically



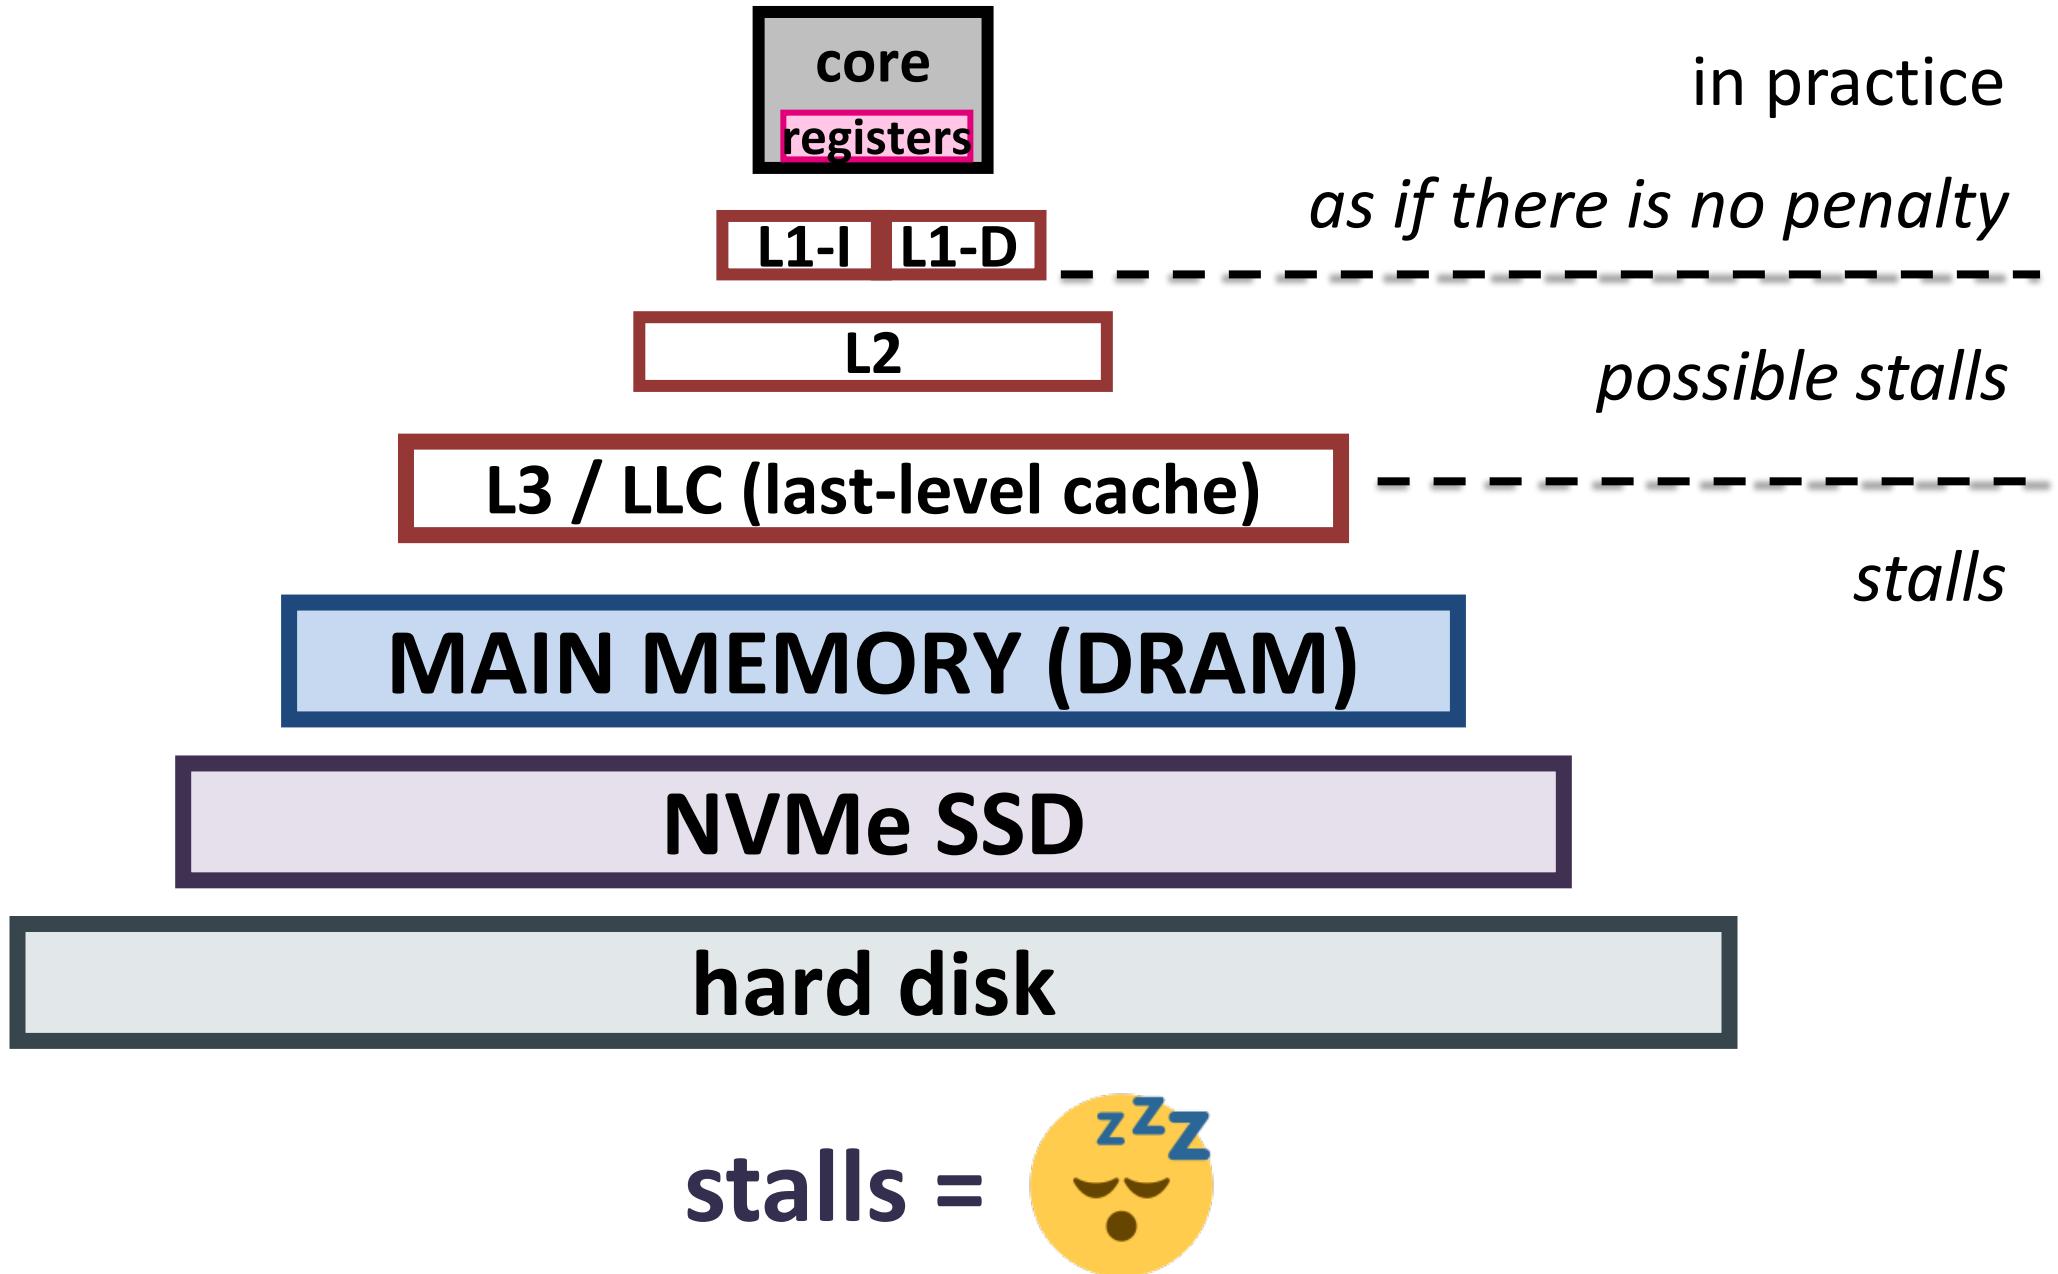
multithreading
threads share
execution cycles
on the same core

why?

we don't want cores to
stay idle waiting for
instruction/data accesses!

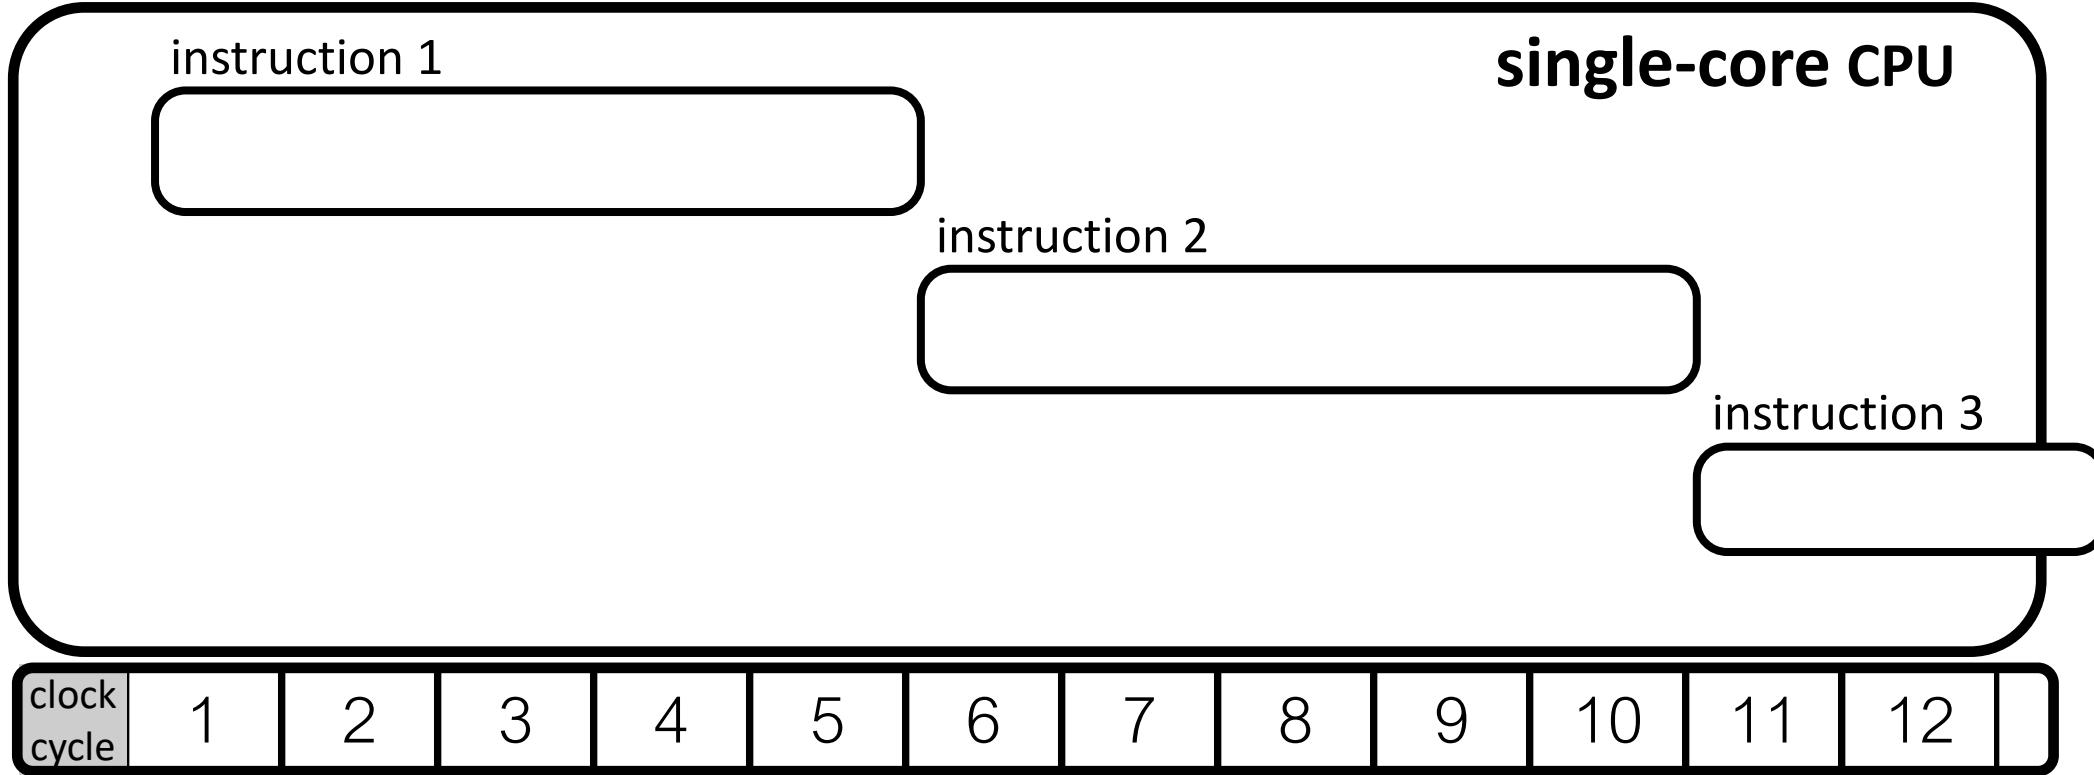
**goal: minimize stall time due to cache/memory accesses
overlapping access latency for one item with other work**

(typical) storage hierarchy – access latency



subscalar CPU (i.e., no implicit parallelism)

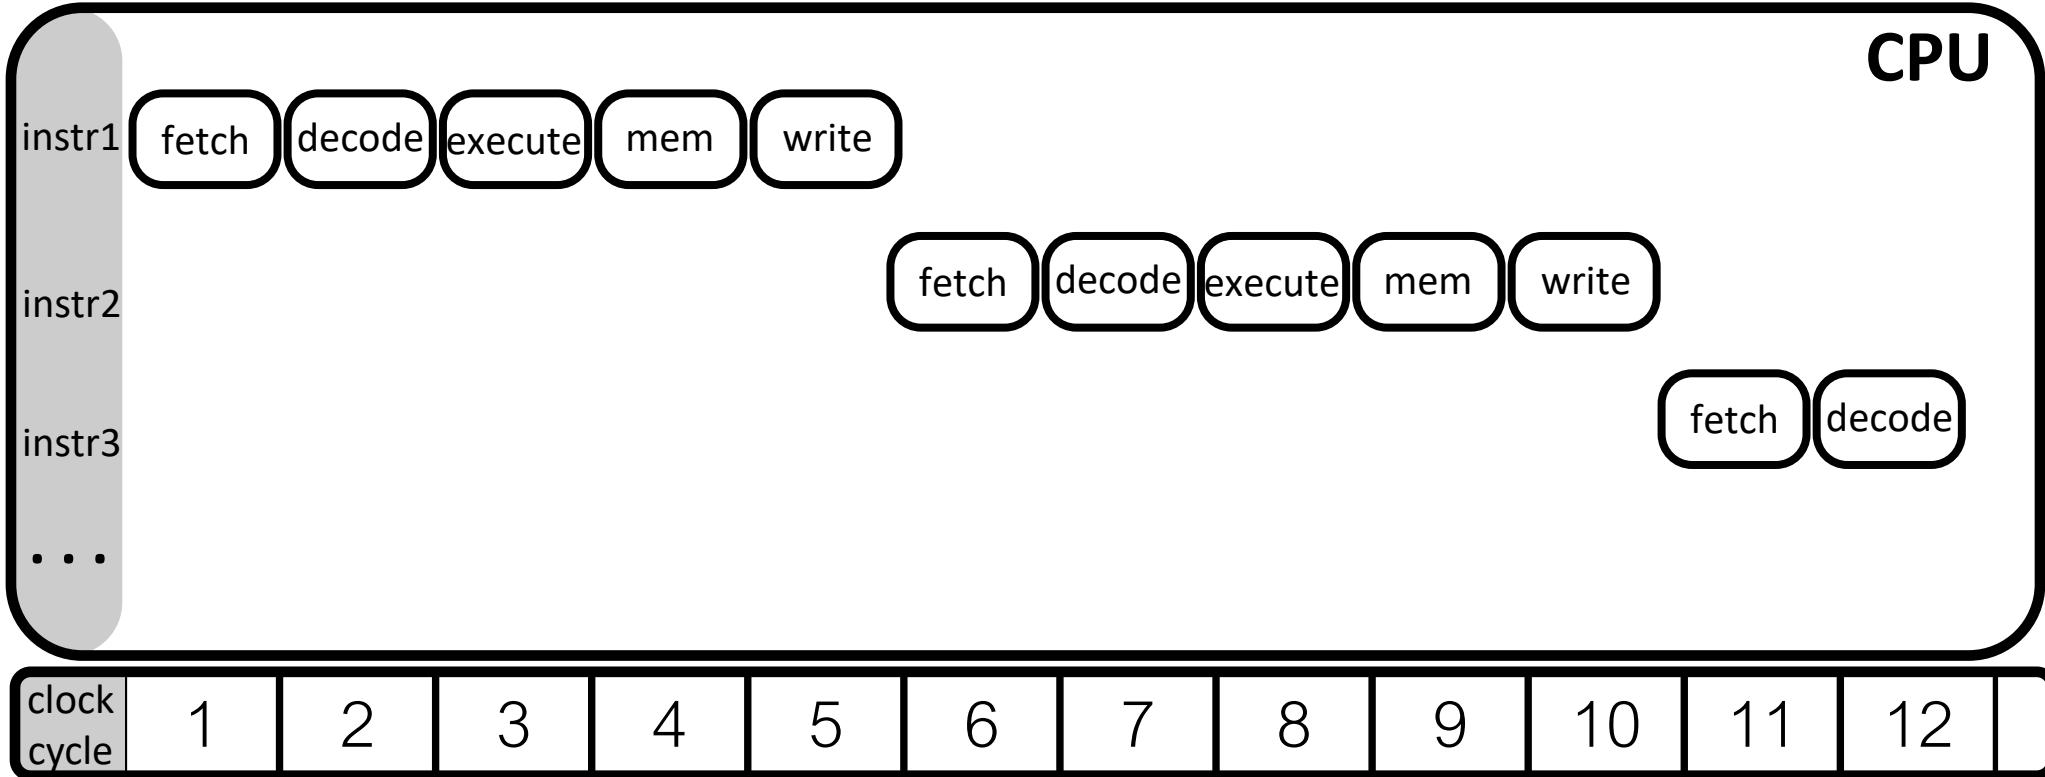
one instruction at a time



several cycles to complete two instructions,
assuming no long-latency data/memory accesses

subscalar CPU (i.e., no implicit parallelism)

one instruction at a time



several cycles to complete two instructions,
assuming no long-latency data/memory accesses

RISC instruction stages

fetch

the instruction from the cache

decode

specifying which operation the instruction performs
and inputs it needs

execute

the operation itself

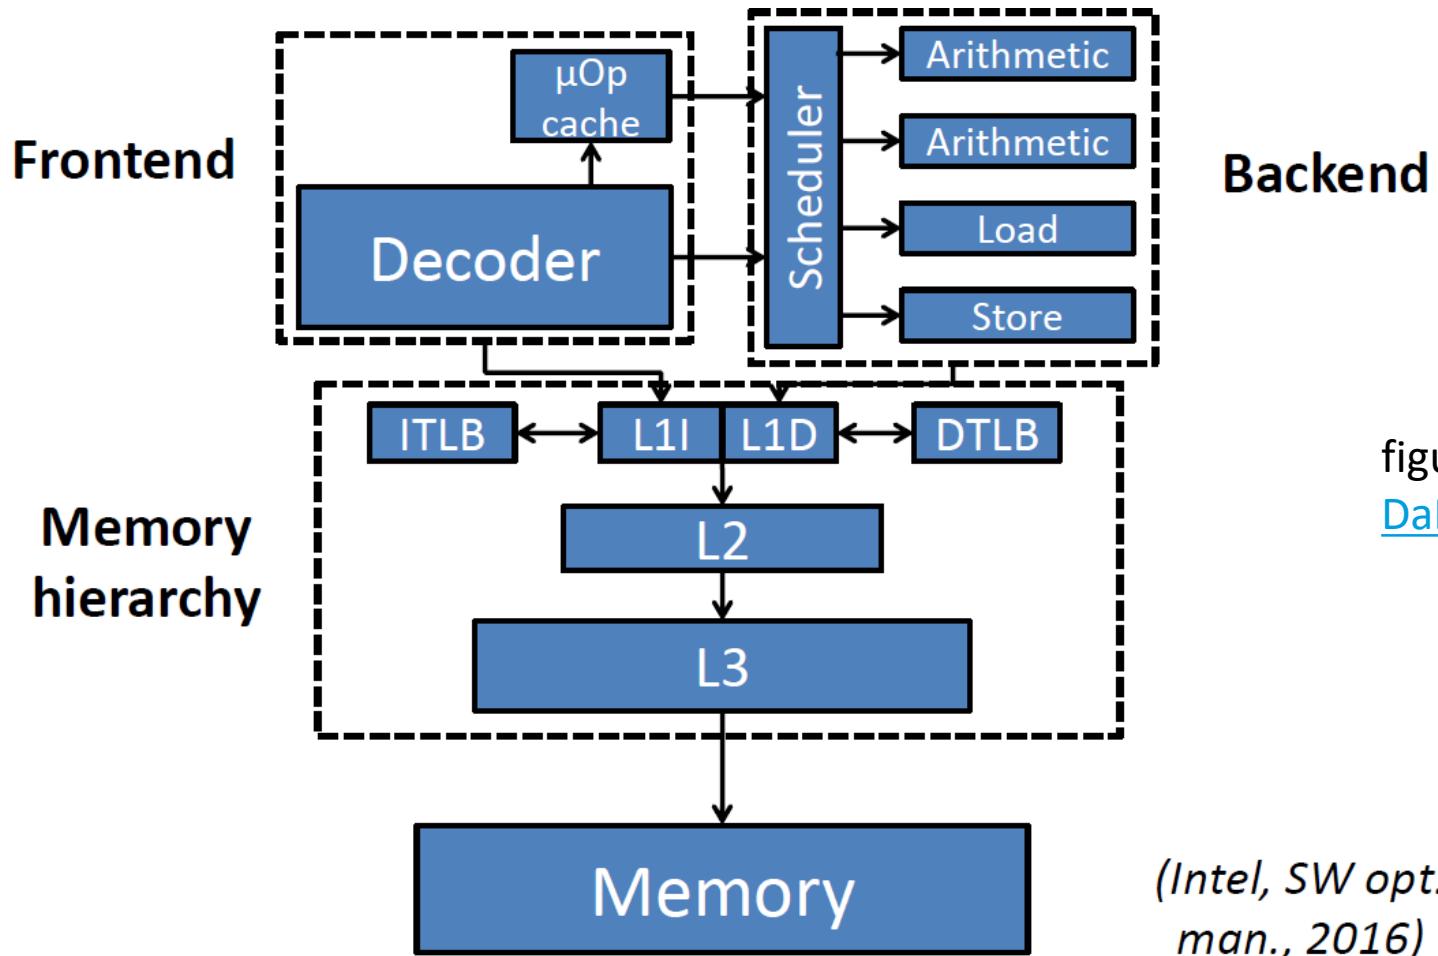
memory

accessing the memory for inputs if needed

write

writing back the results into registers

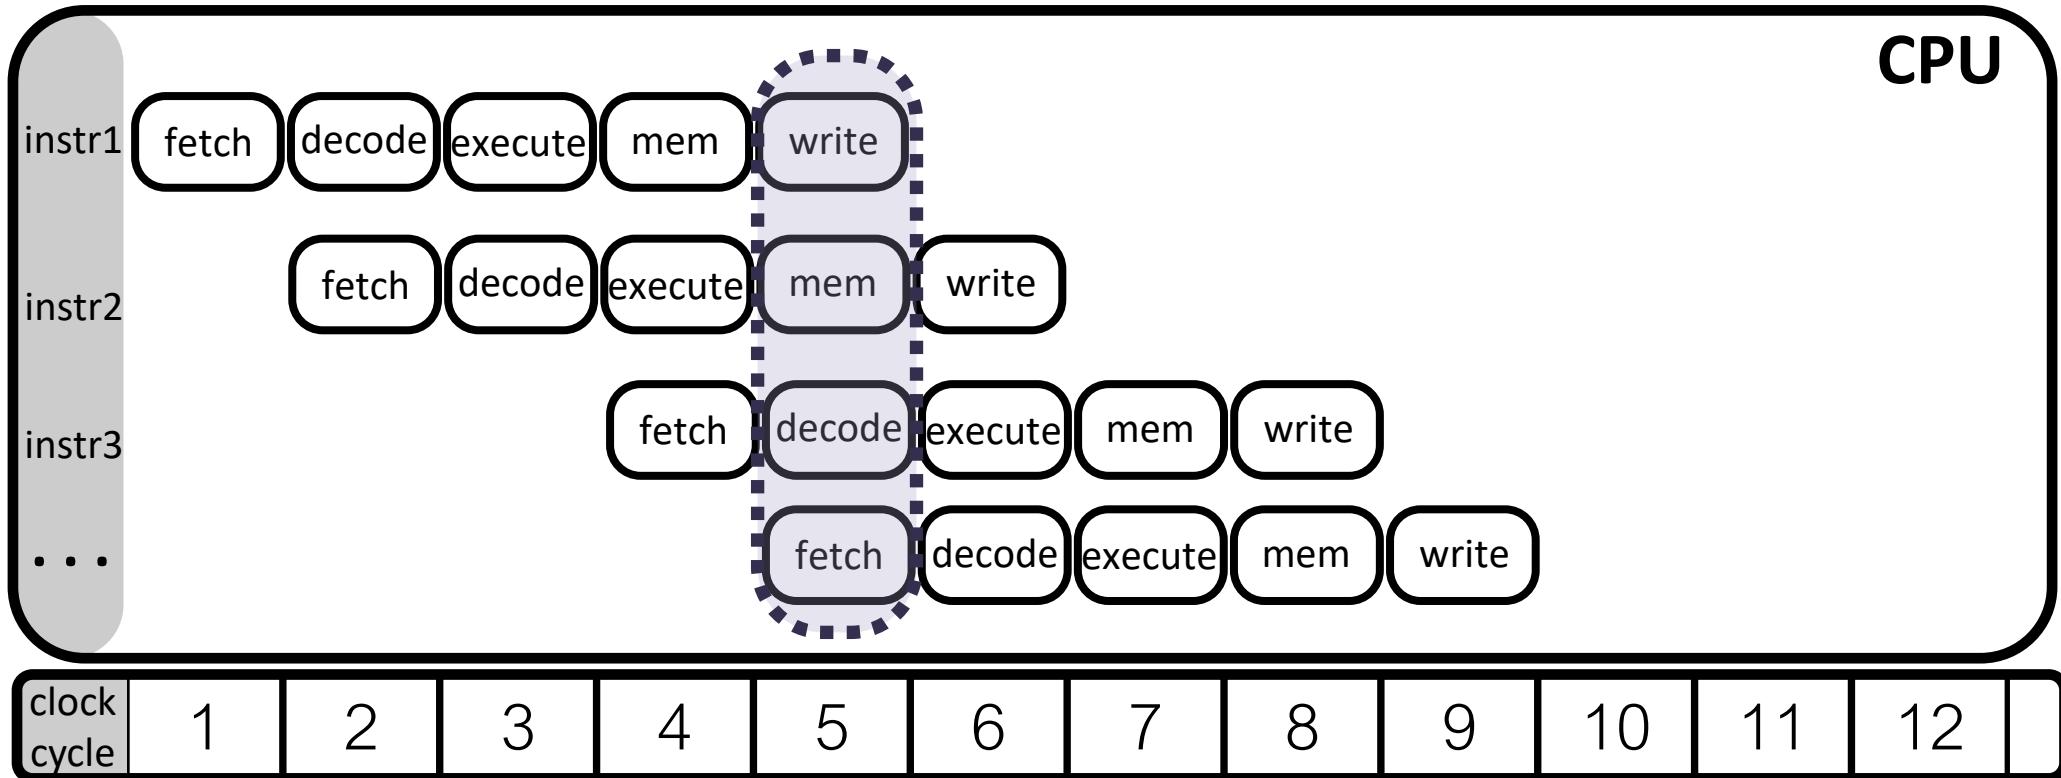
micro-architecture of an OoO processor



delays in fetching, decoding, etc. an instruction
cause *frontend stalls*, rest cause *backend stalls*

instruction pipelining

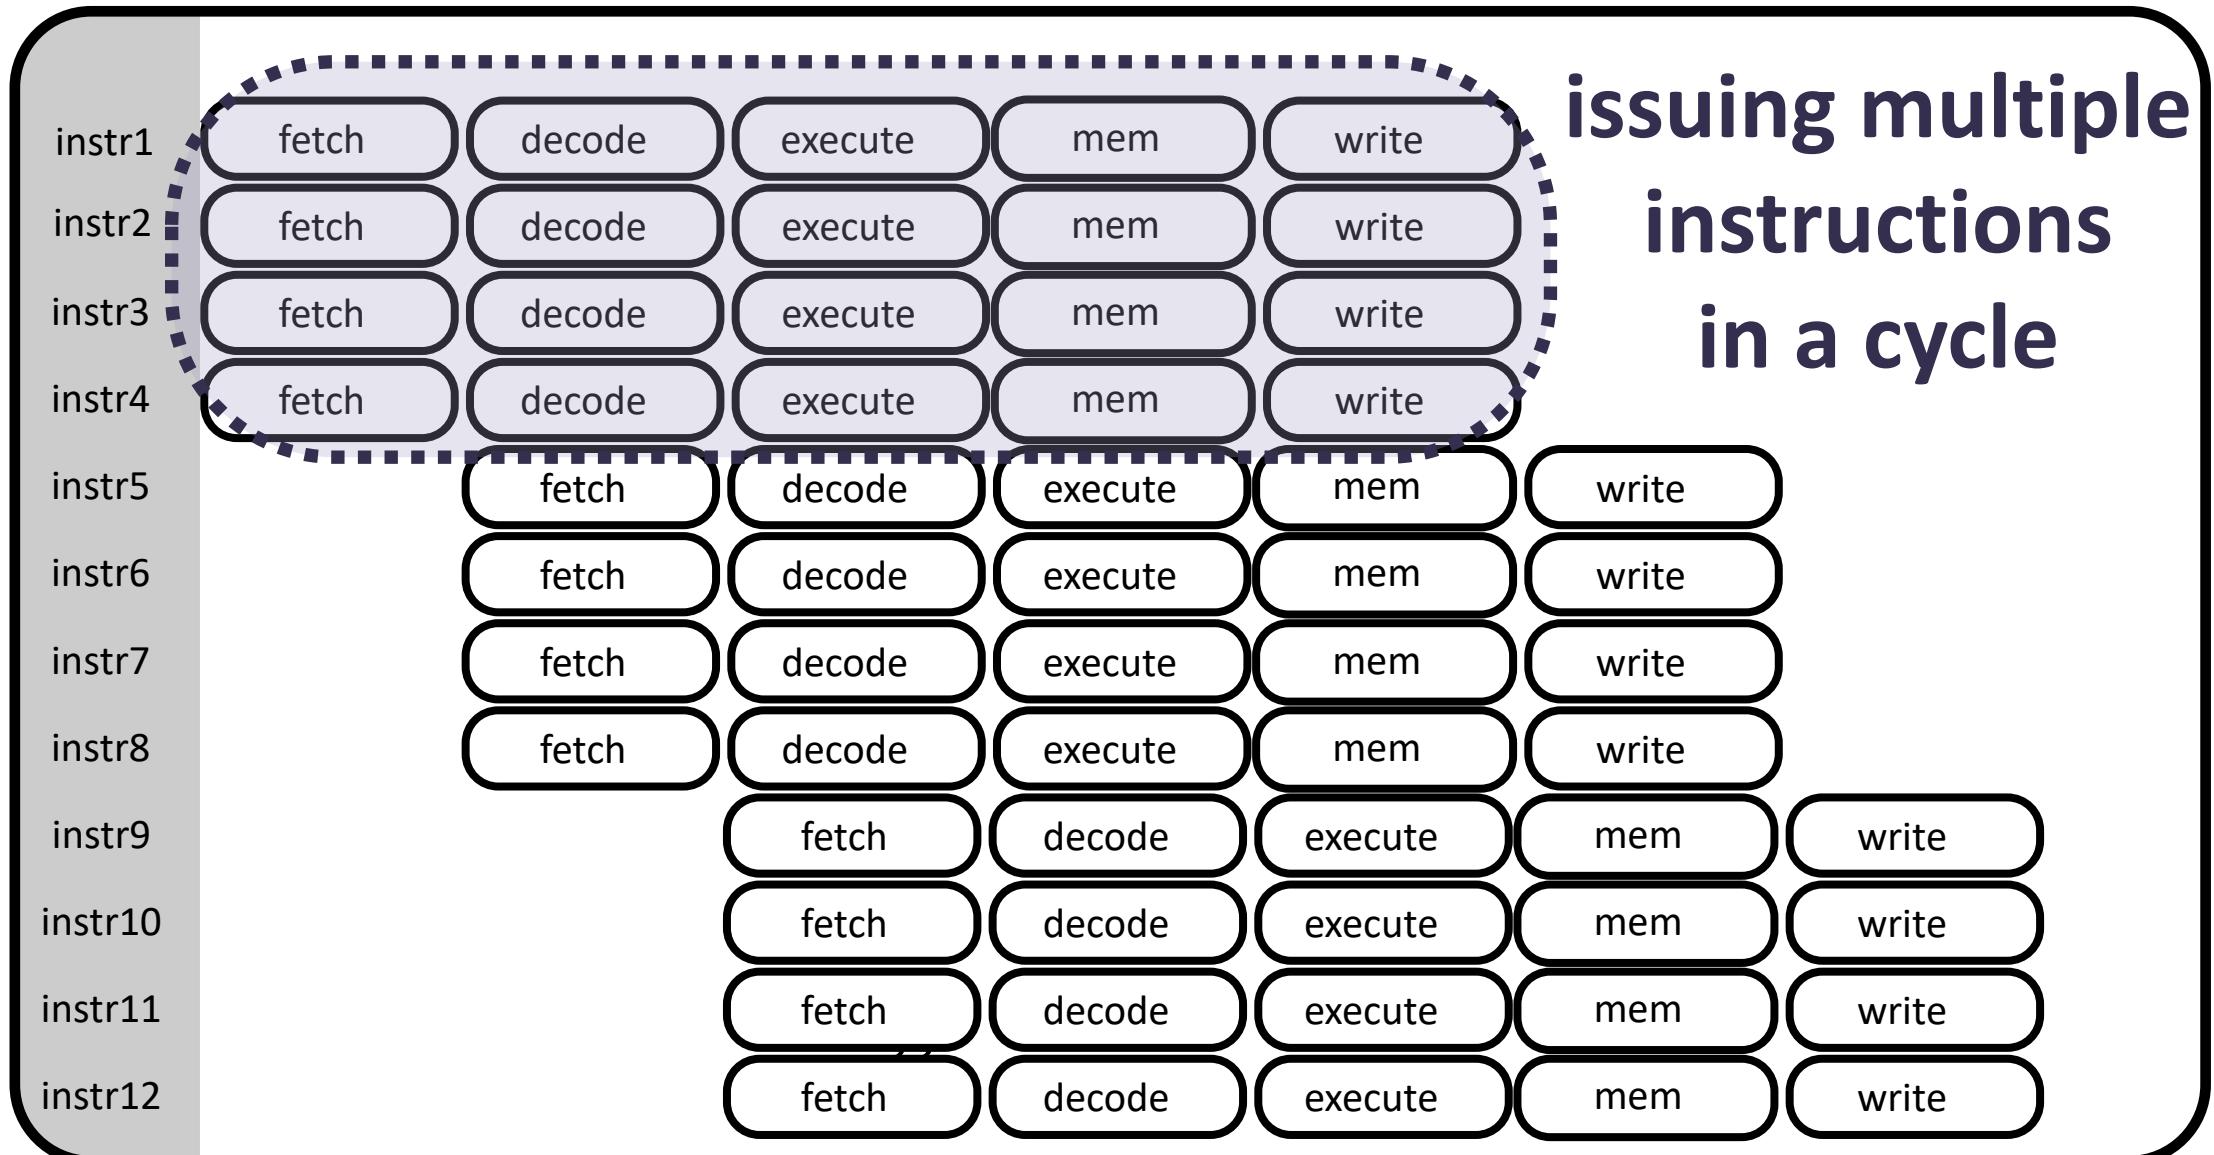
overlapping stages of different instructions



fundamental way to parallelize implicitly

superscalar CPU

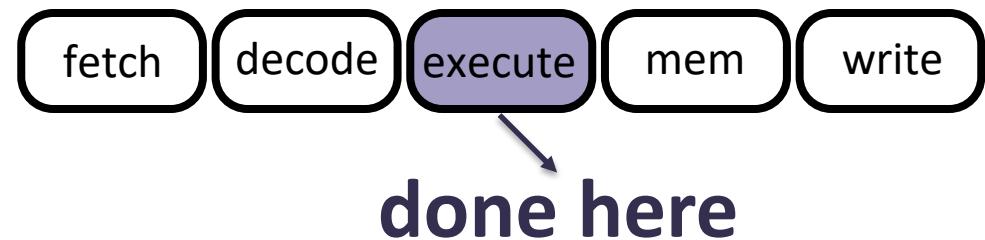
example of 4-way superscalar CPU



issuing multiple
instructions
in a cycle

out-of-order (OoO) execution

allows a processor to execute instructions based on the availability of input data rather than strictly following the instruction ordering of a program

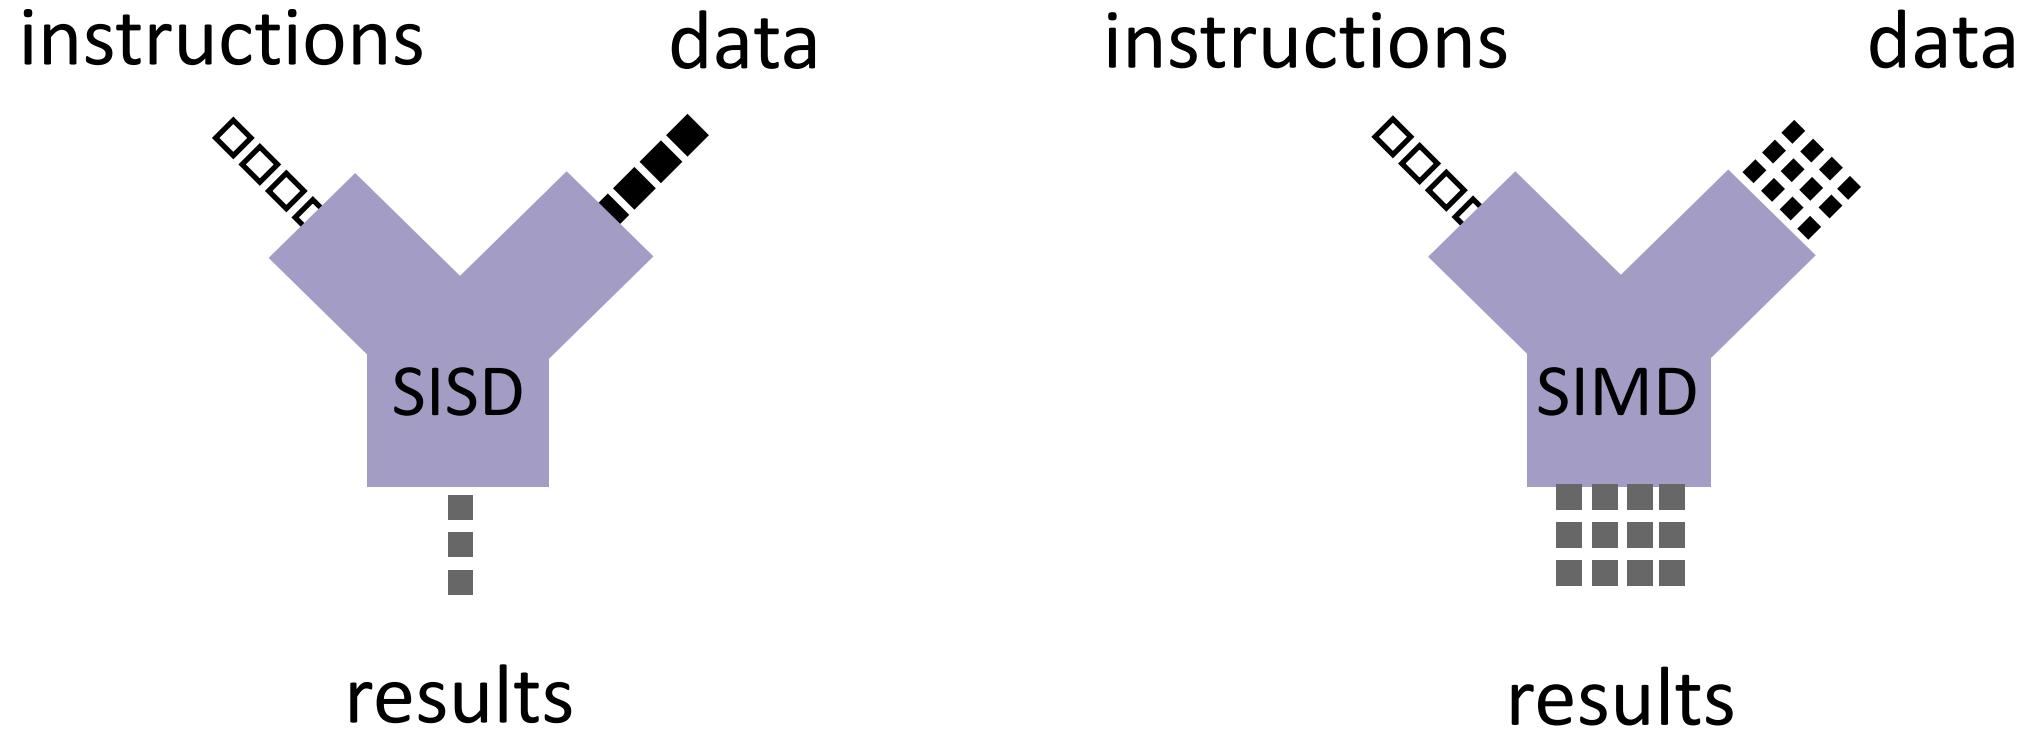


example:

- (1) $r1 \leftarrow r2 / r3$
- (2) $r4 \leftarrow r1 + r5$
- (3) $r6 \leftarrow r7 * r8$

independent from the previous two,
can be executed independently in
parallel or before 1 & 2

single-instruction multiple data (SIMD)



**also, implicit data parallelism
but this one has to be managed/issued by the software**

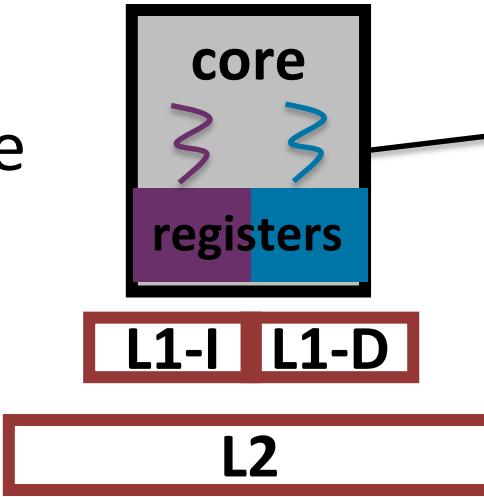
simultaneous multithreading (SMT)

keep state/context of multiple threads via more register space

each CPU cycle,
the threads are swapped

no need for a context switch
for this by the operating system

today, also known as hyperthreading



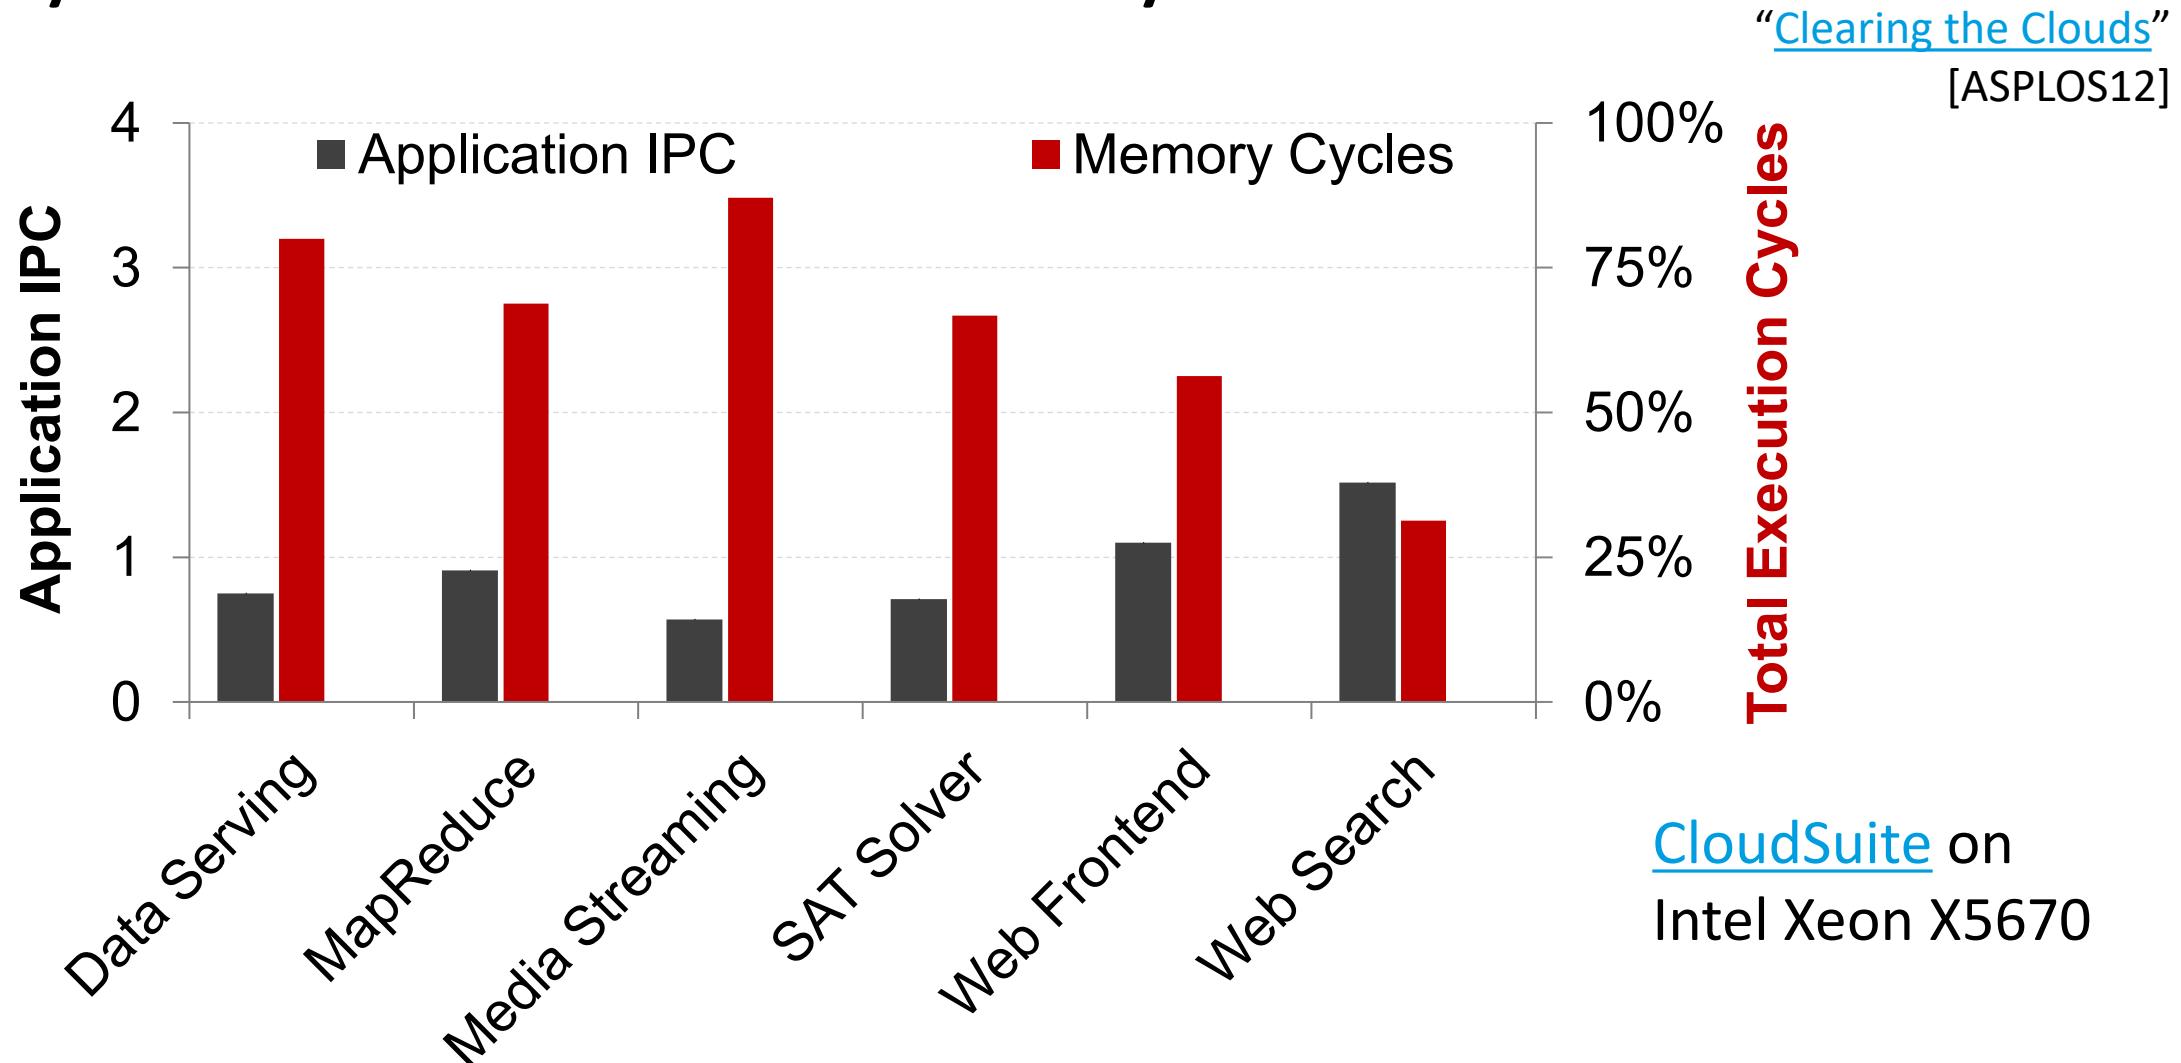
a core that supports
two hardware contexts
(or logical threads)

design trade-off: how to use
the transistors in a core?

- to create higher per-thread speed by making the instruction pipeline more complex
- vs more threads

**more pressure on shared hardware resources (e.g., caches)
if not used right, may even hurt performance**

memory stalls in data-intensive systems



data-intensive systems are known to suffer from memory stalls

agenda

lecture

- storage hierarchy overview
- local memory hierarchy: DRAM & caches
- hardware parallelism: implicit & explicit
- different memory hierarchies

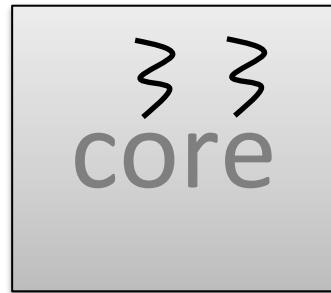
exercises

- paper discussion
- c/cpp examples
- setting core affinity

types of hardware parallelism

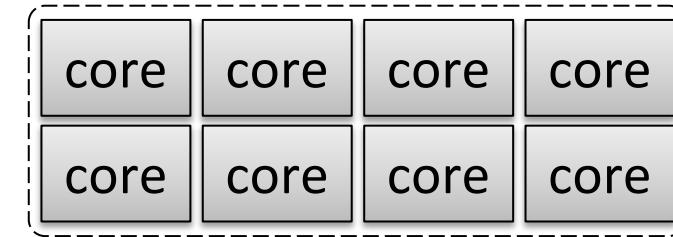
implicit/vertical parallelism

explicit/horizontal parallelism



single-core

instruction & data parallelism
simultaneous multithreading



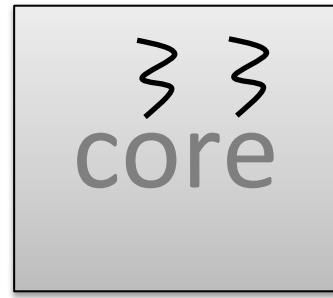
multicores

multiple threads run
in parallel on different cores

implicit parallelism → (almost) free lunch
explicit parallelism → must work hard to leverage it

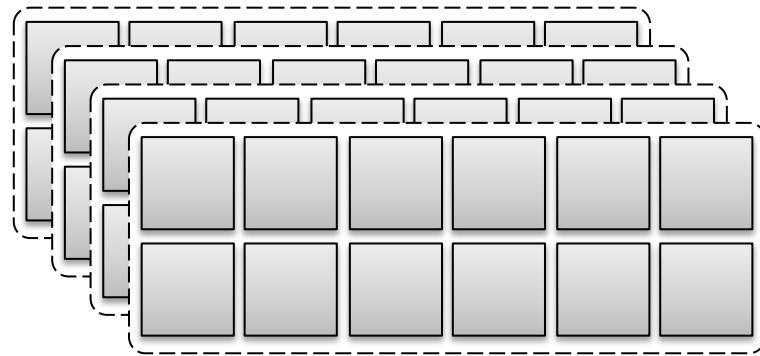
types of hardware parallelism

implicit/vertical parallelism



single-core
instruction & data parallelism
simultaneous multithreading

explicit/horizontal parallelism

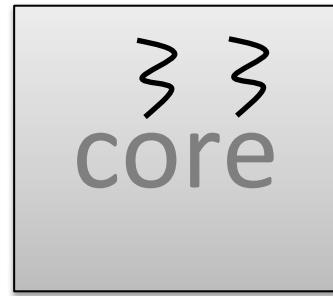


multisocket multicores
multiple processors/ CPUs
in one machine

implicit parallelism → (almost) free lunch
explicit parallelism → must work hard to leverage it

types of hardware parallelism

implicit/vertical parallelism



explicit/horizontal parallelism



single-core

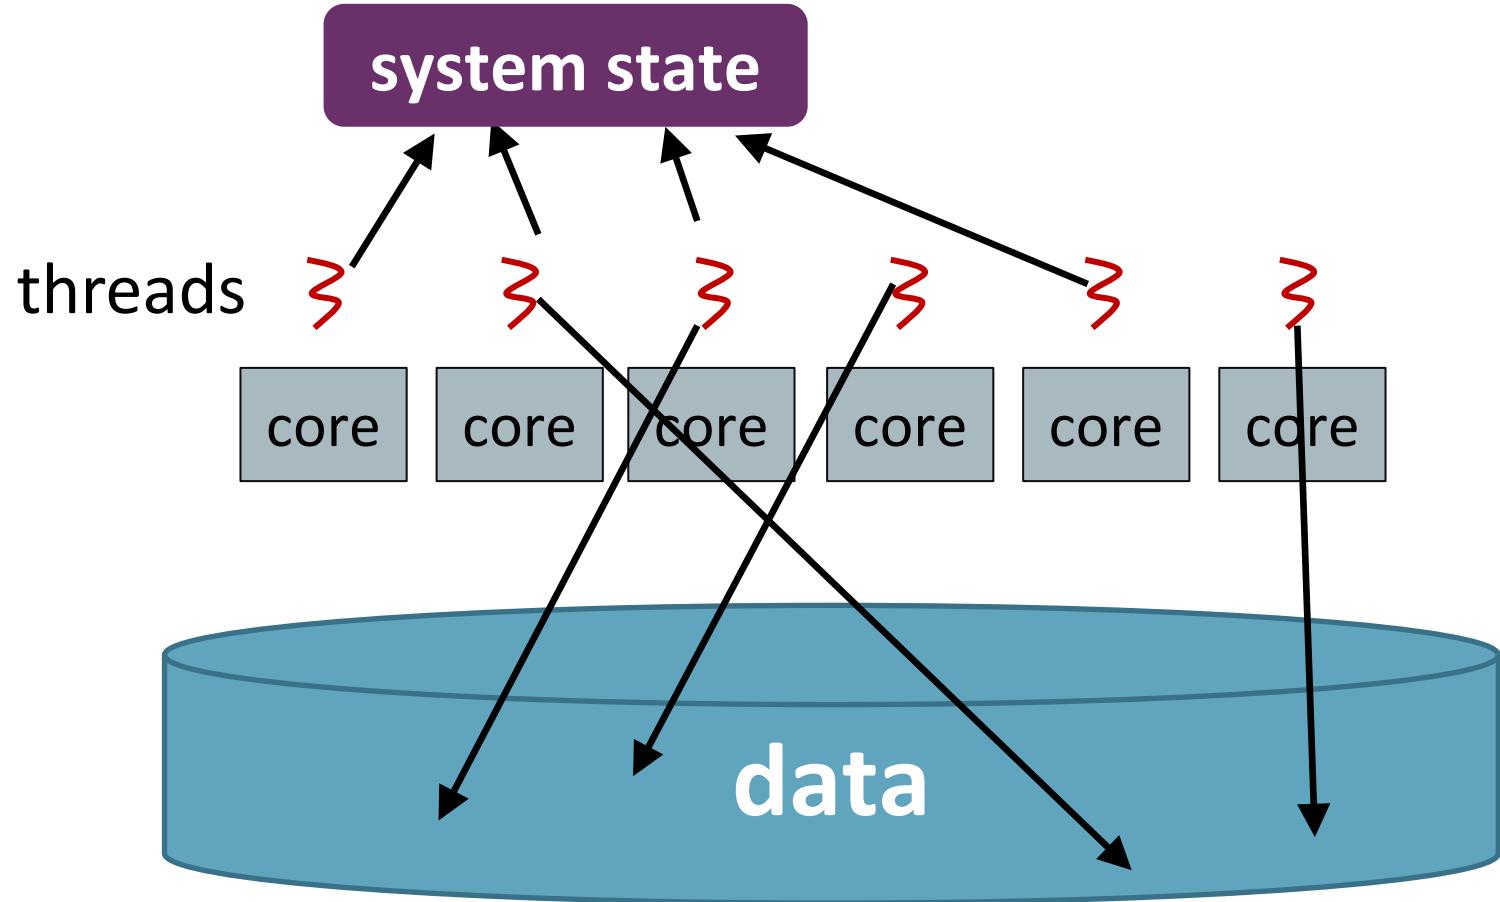
instruction & data parallelism
simultaneous multithreading

distributed systems
running a program over
multiple machines

implicit parallelism → (almost) free lunch

explicit parallelism → must work hard to leverage it

why explicit parallelism is hard to leverage?



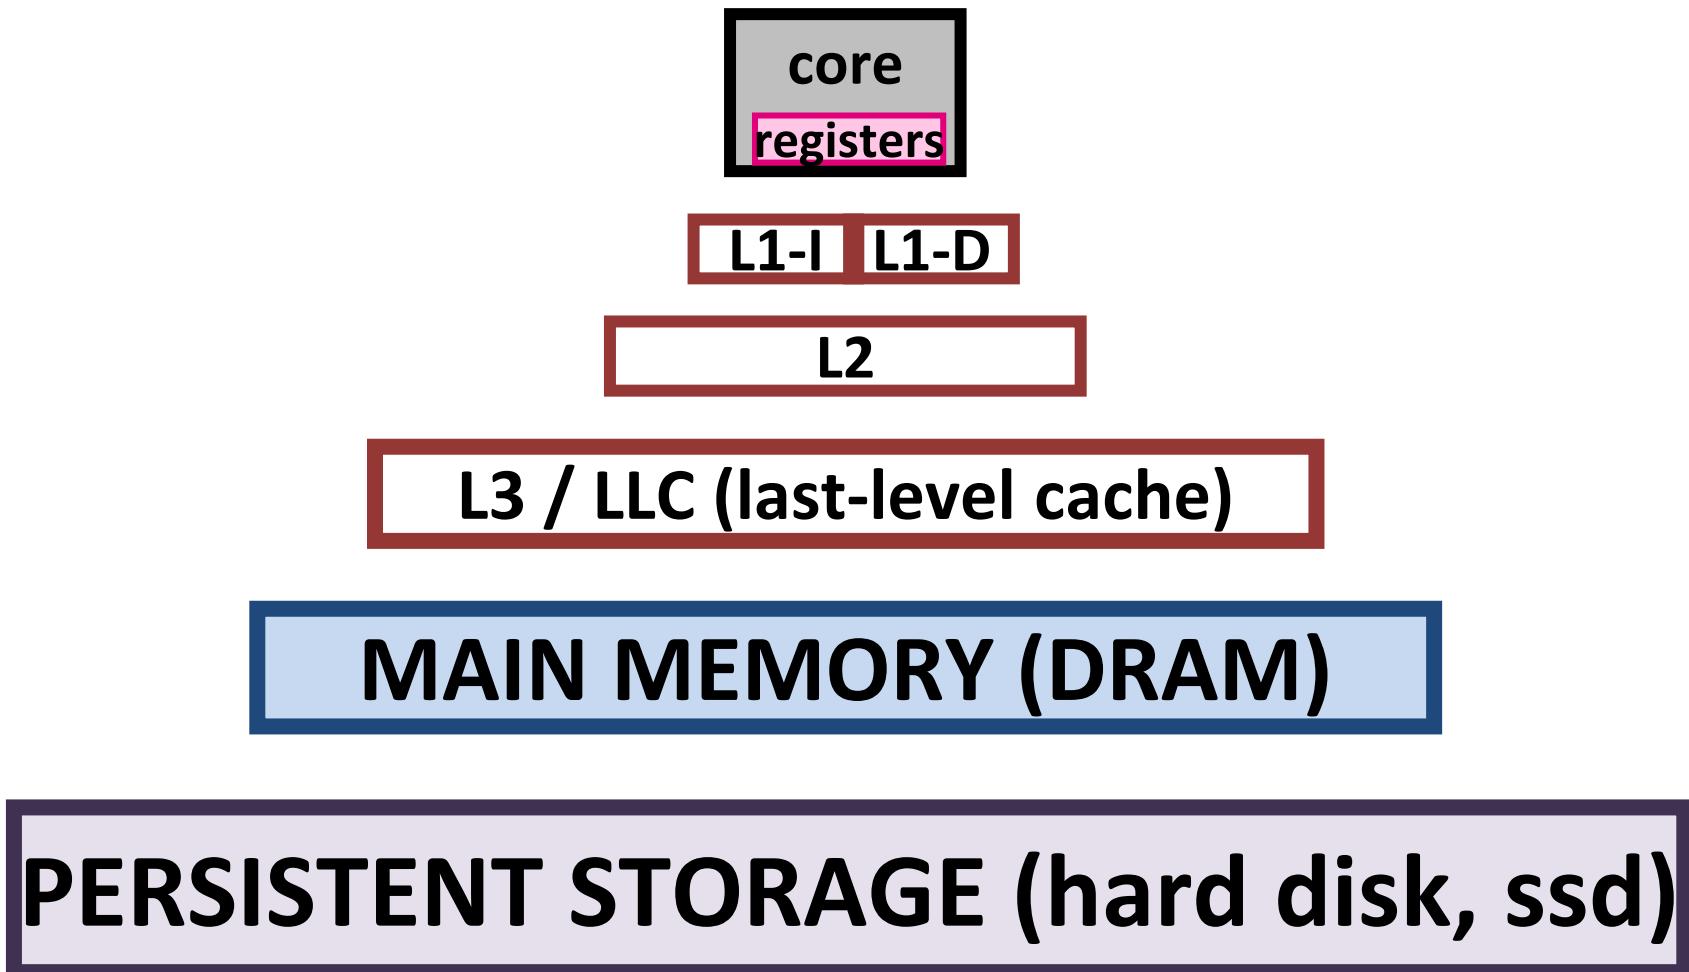
synchronization



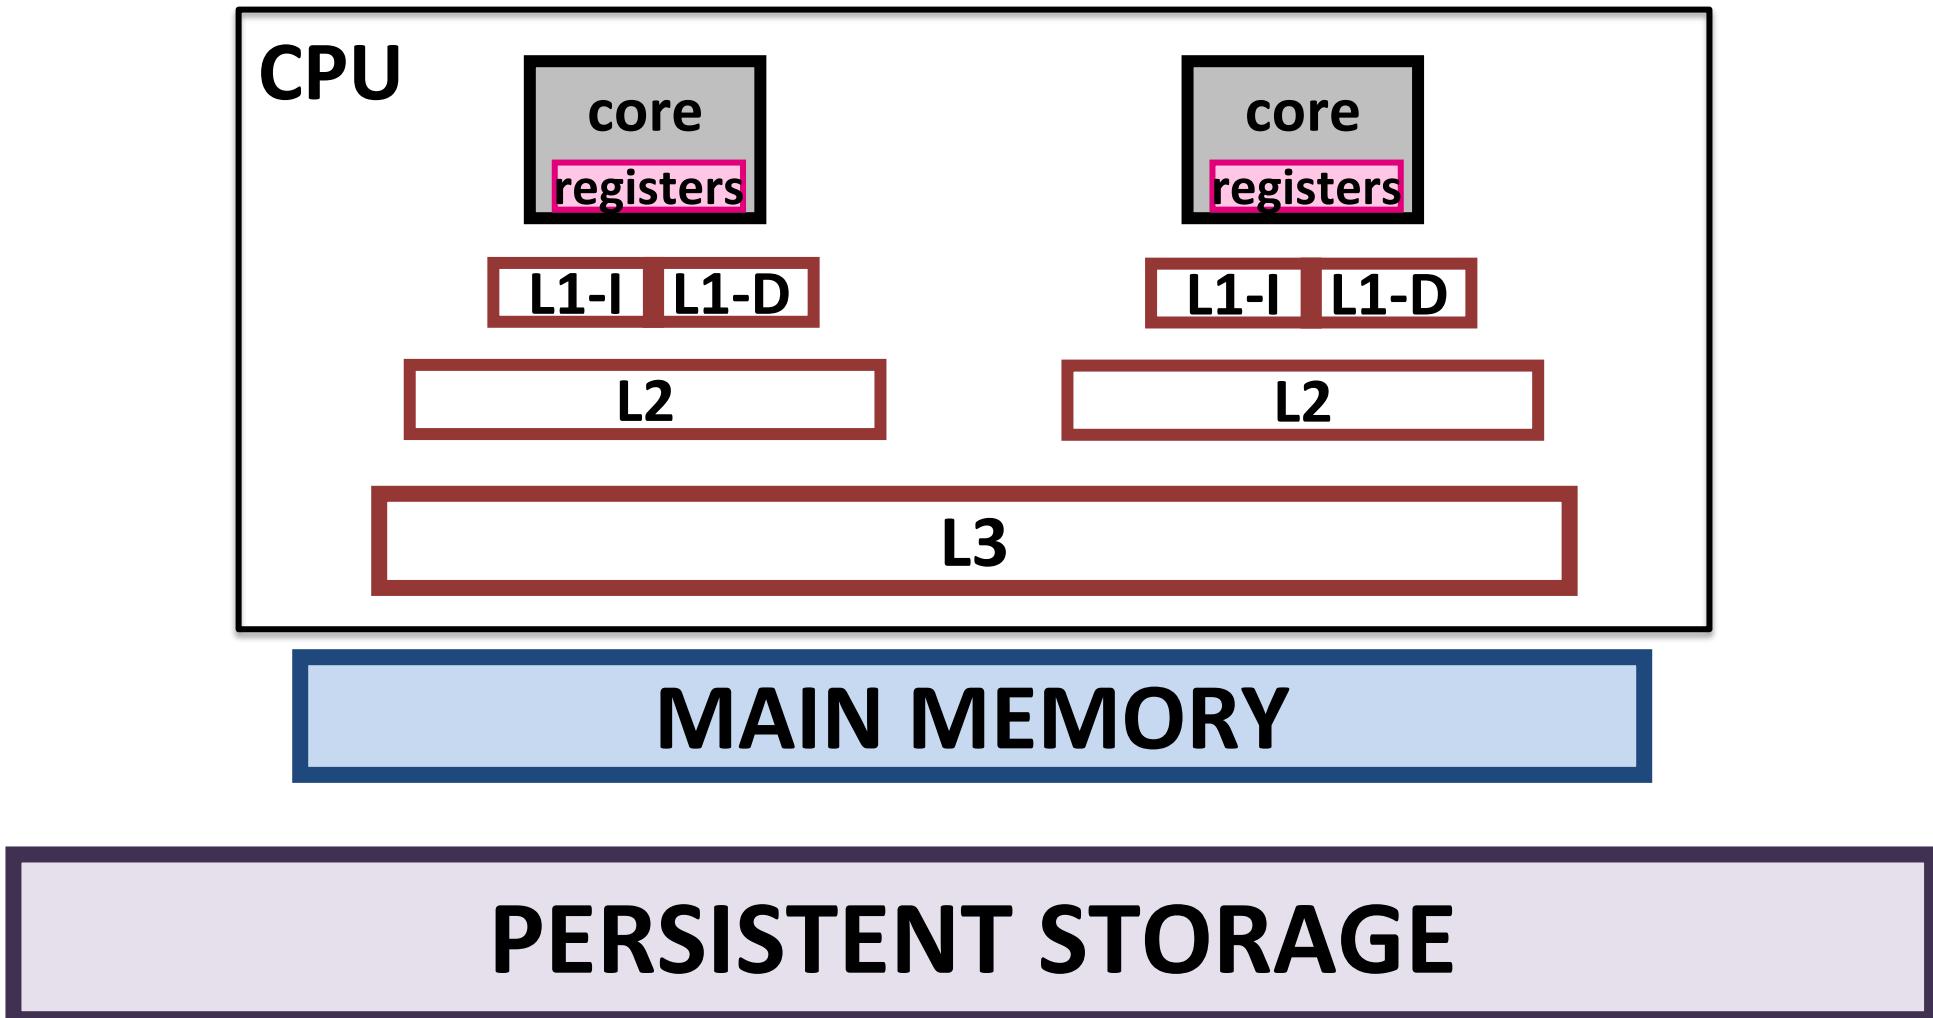
**too many locks →
low performance**

**if a system has many unpredictable accesses to shared data,
need to synchronize the threads while accessing to data**

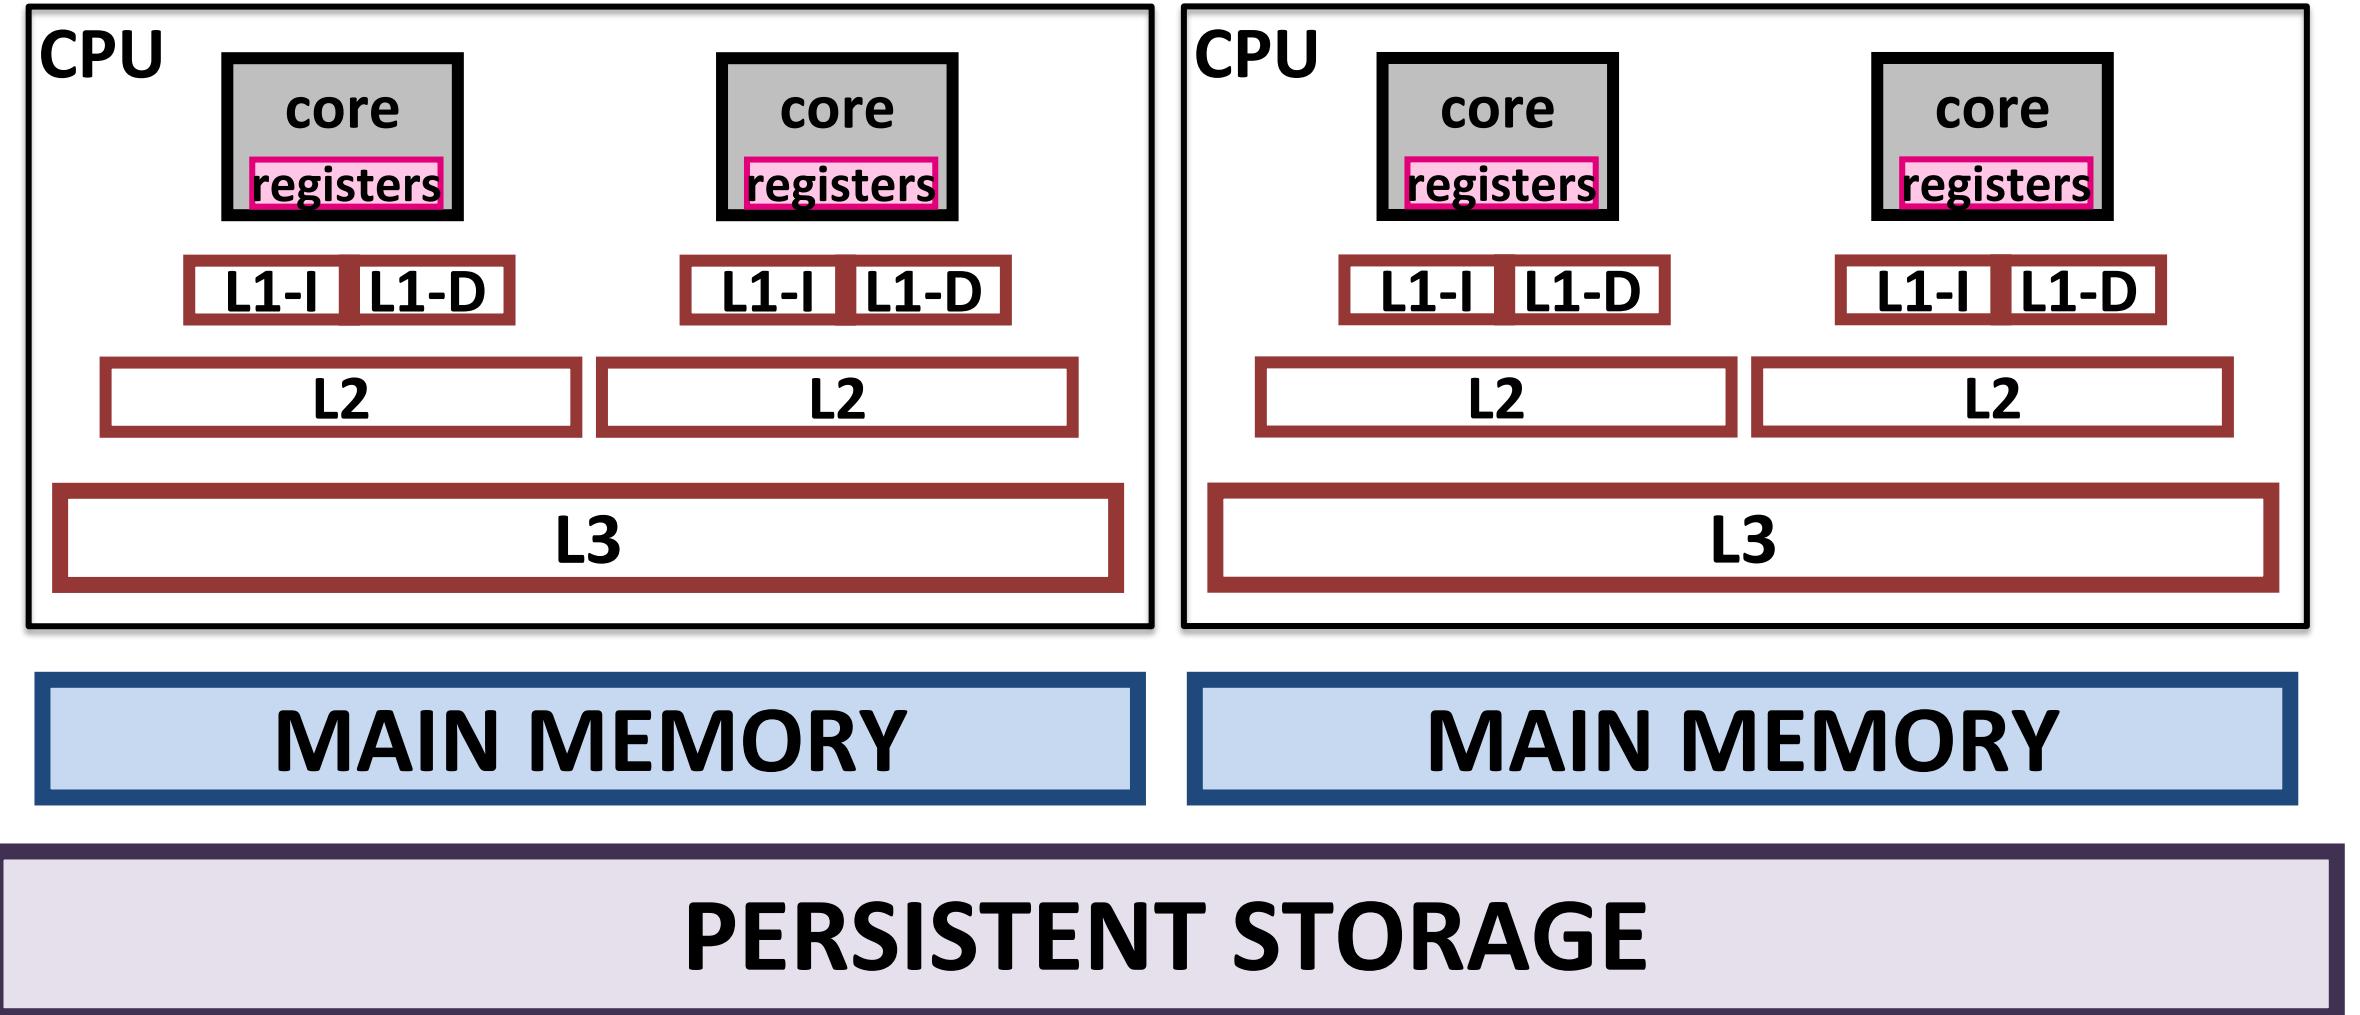
single-core CPU storage hierarchy



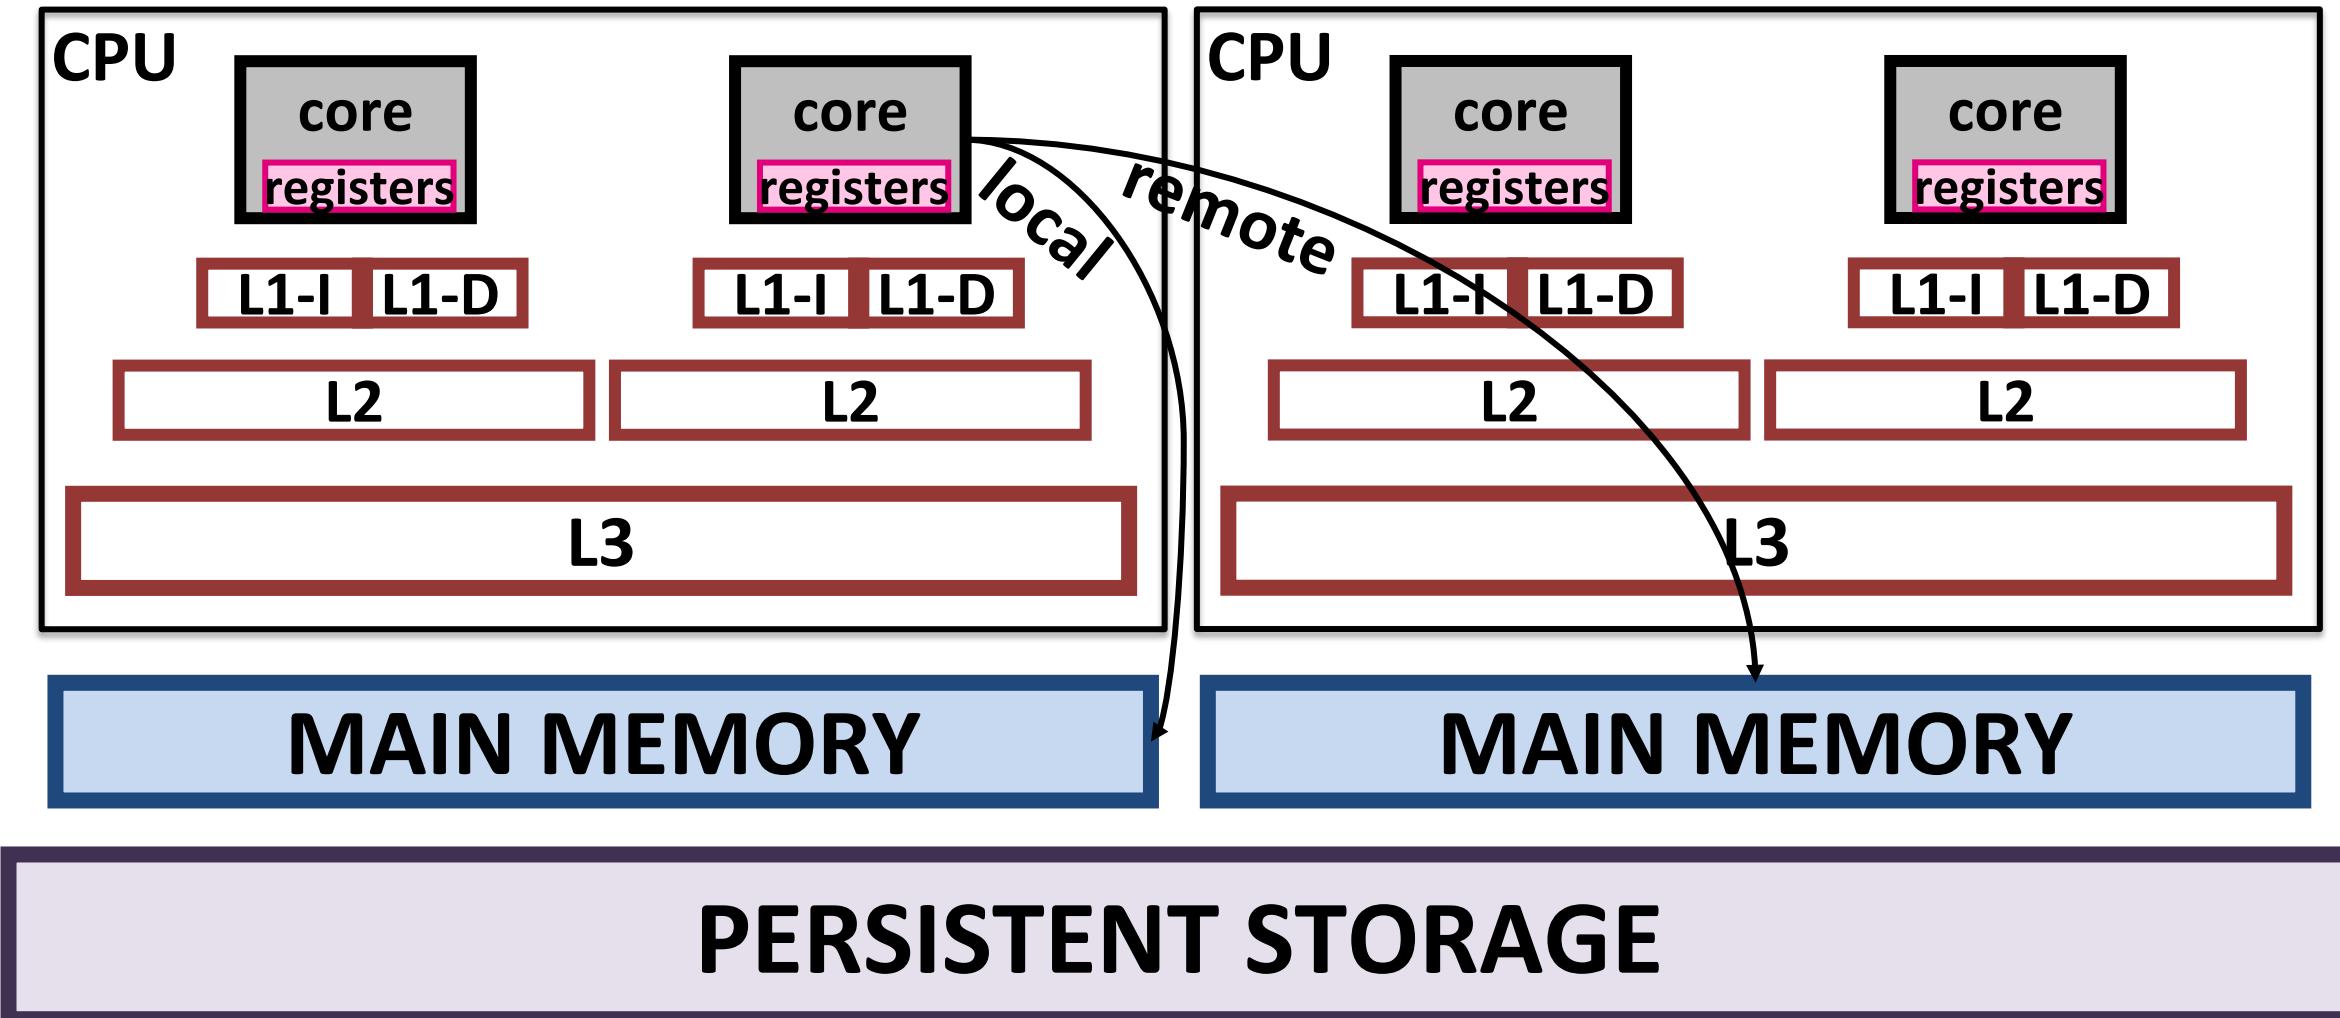
multicore CPU storage hierarchy



multi-socket multicore CPU storage hierarchy

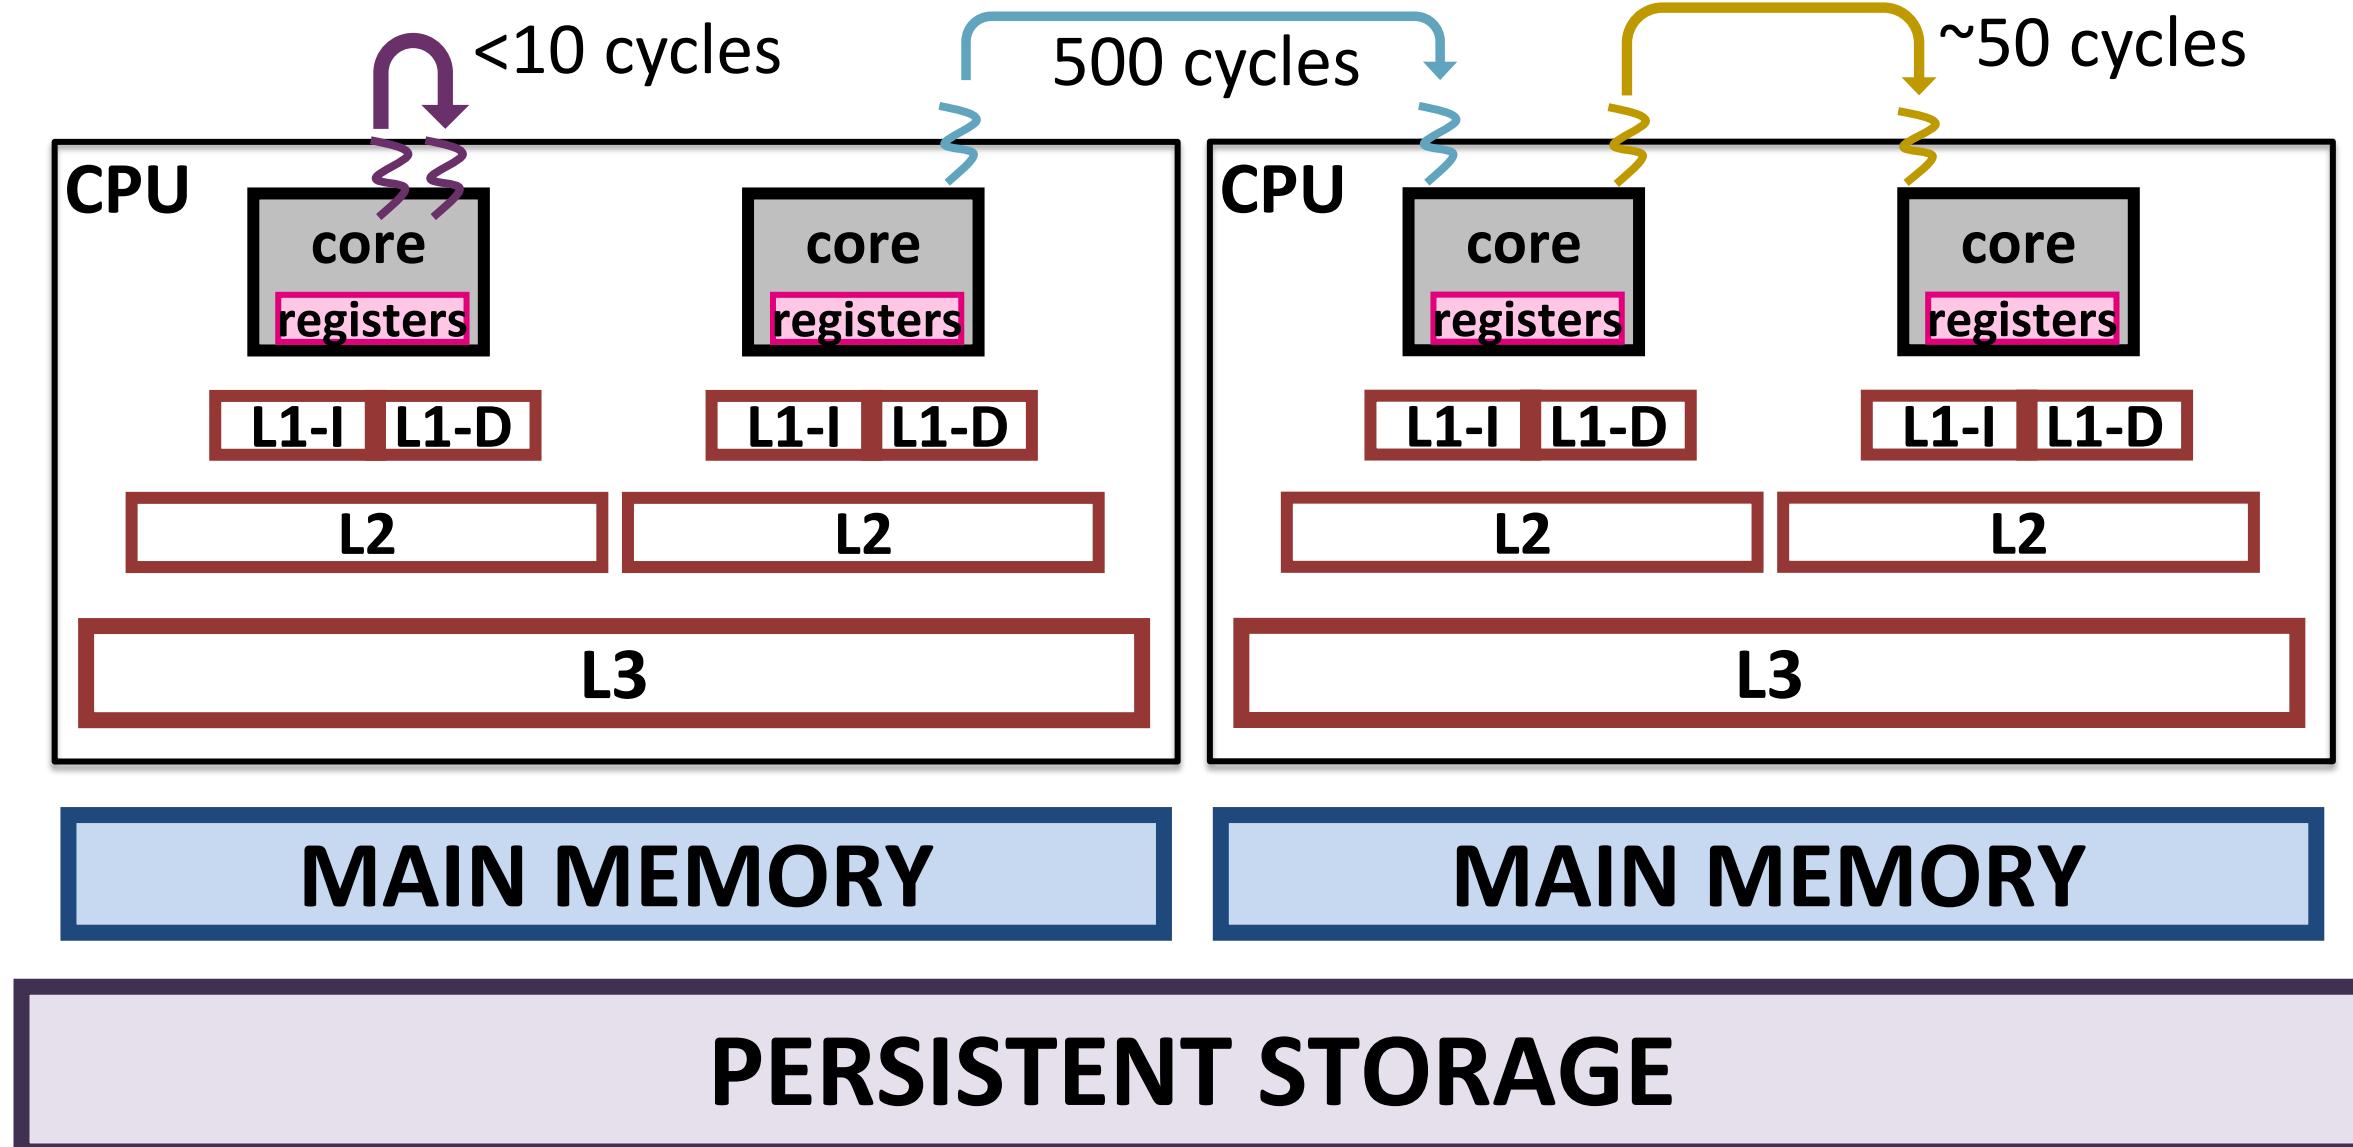


NUMA = non-uniform memory access



local memory access is faster than remote one!

NUMA impact



agenda

lecture

- storage hierarchy overview
- local memory hierarchy: DRAM & caches
- hardware parallelism: implicit & explicit
- different memory hierarchies

exercises

- paper discussion
- c/cpp examples
- setting core affinity

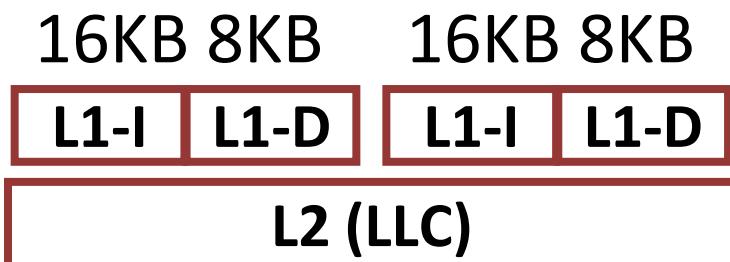
Sun SPARC CPUs

discontinued!

data-intensive workloads have

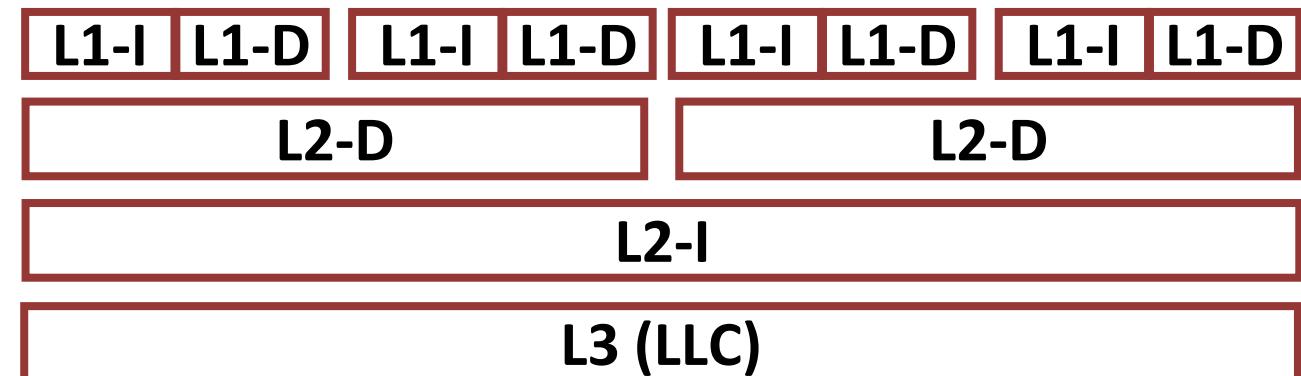
- 1. large instruction footprint**
- 2. low temporal data locality**

UltraSPARC T2 (Niagara 2)



used in “data partitioning
on chip multiprocessors” paper

SPARC M7 & M8



let's check some hardware topology!

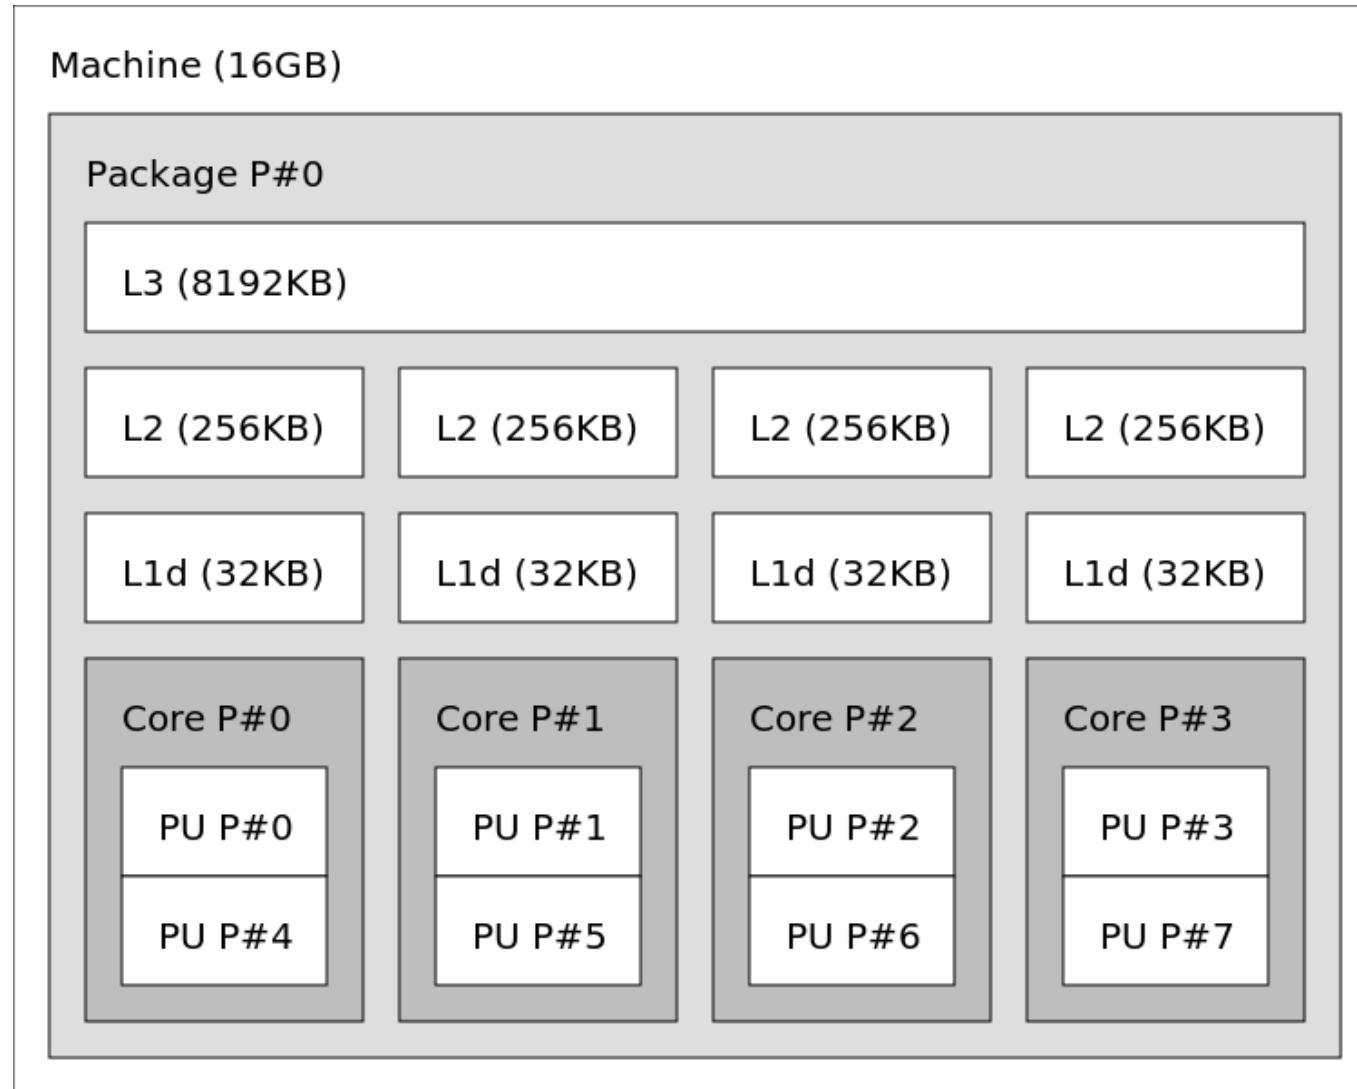
```
$ lstopo
```

or for better visuals

```
$ lstopo --output-format png -v --no-io > cpu.png
```

- try this out on your laptop
- report what you see
- does it look like the hierarchy we have been discussing?
- if not, what is different?

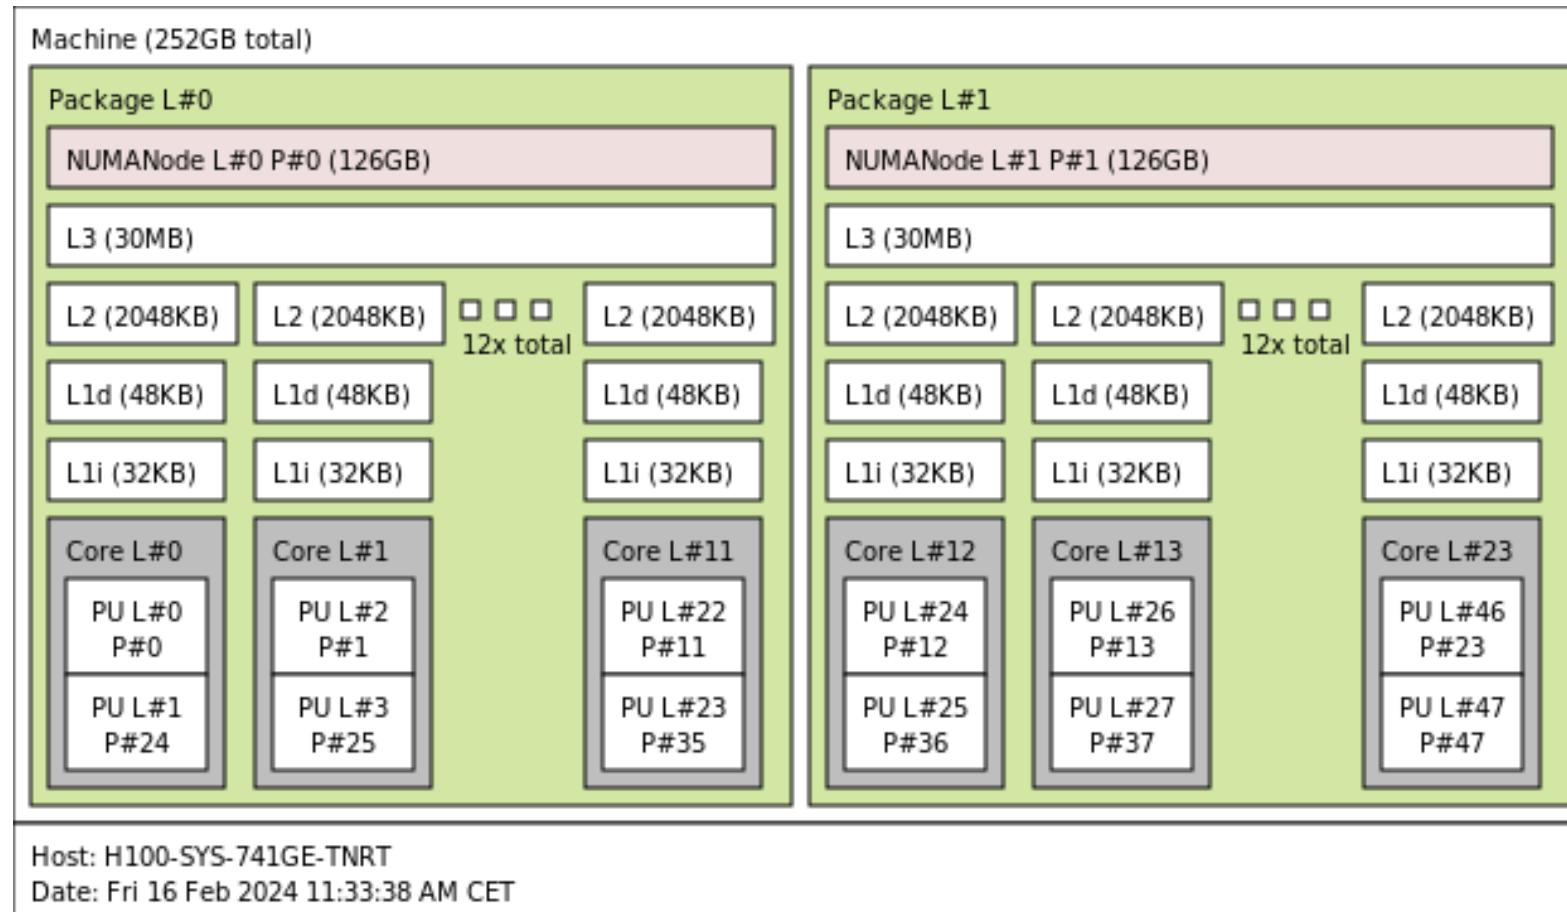
laptop with Intel CPU – hyper-threading on



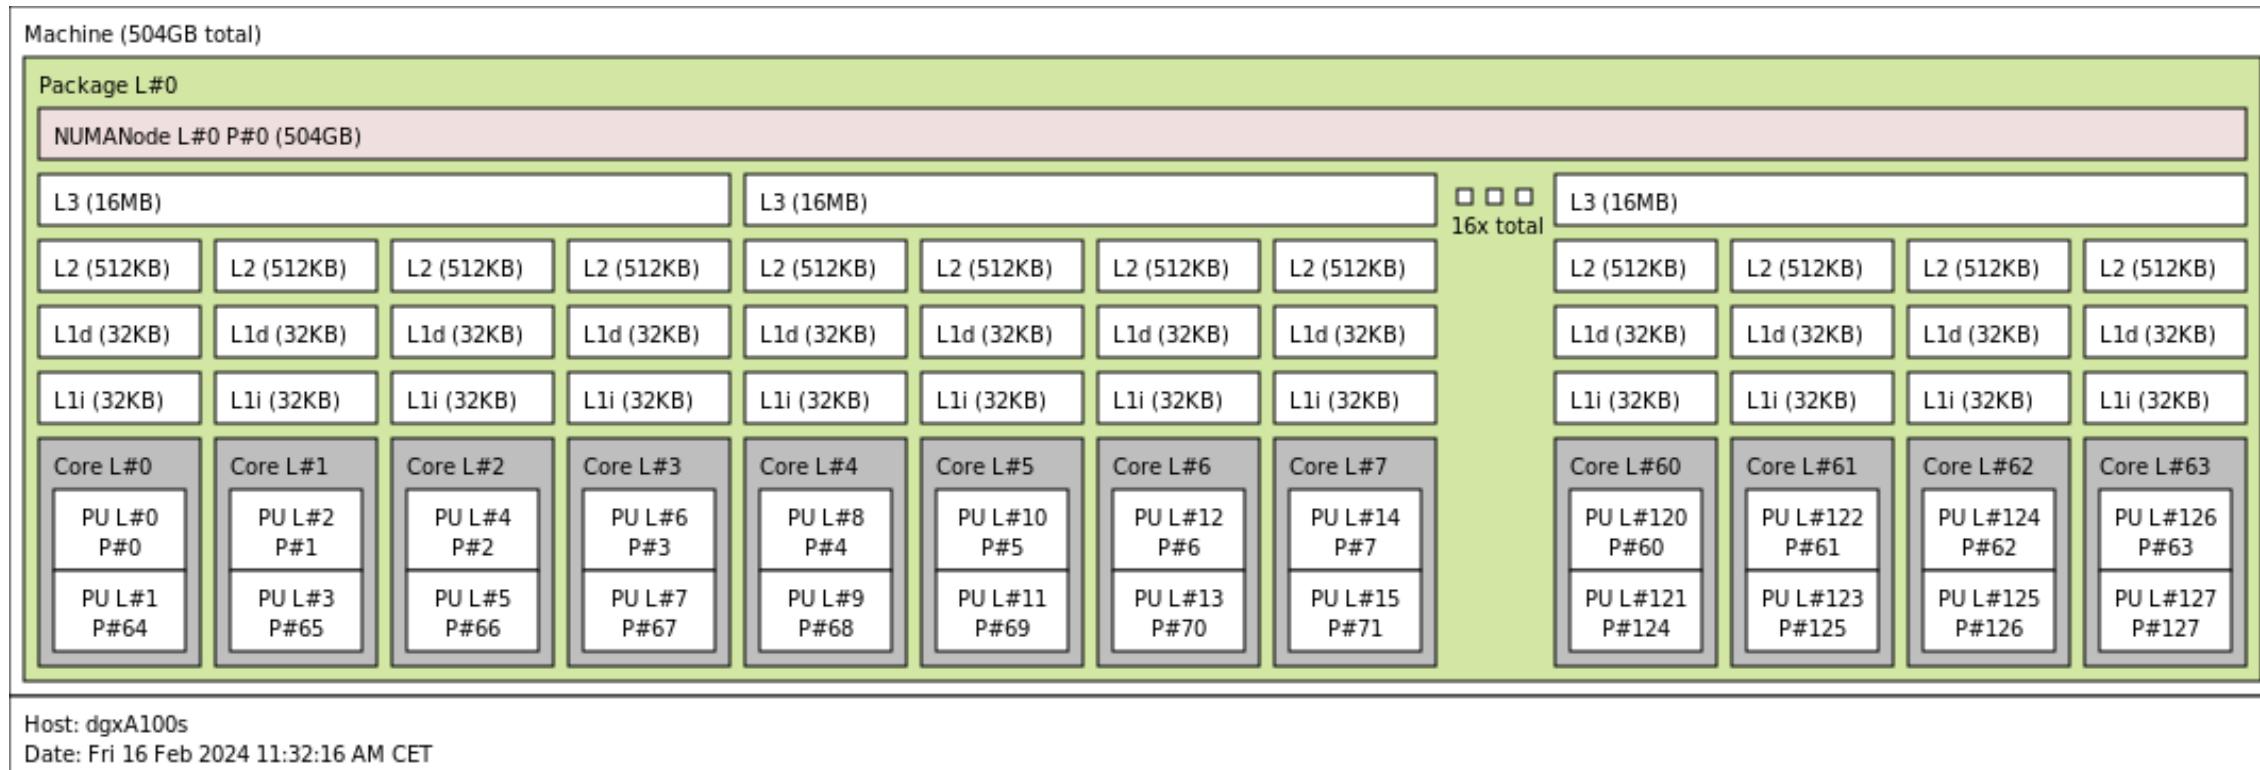
class server – hyper-threading off



newer intel CPUs – hyper-threading on



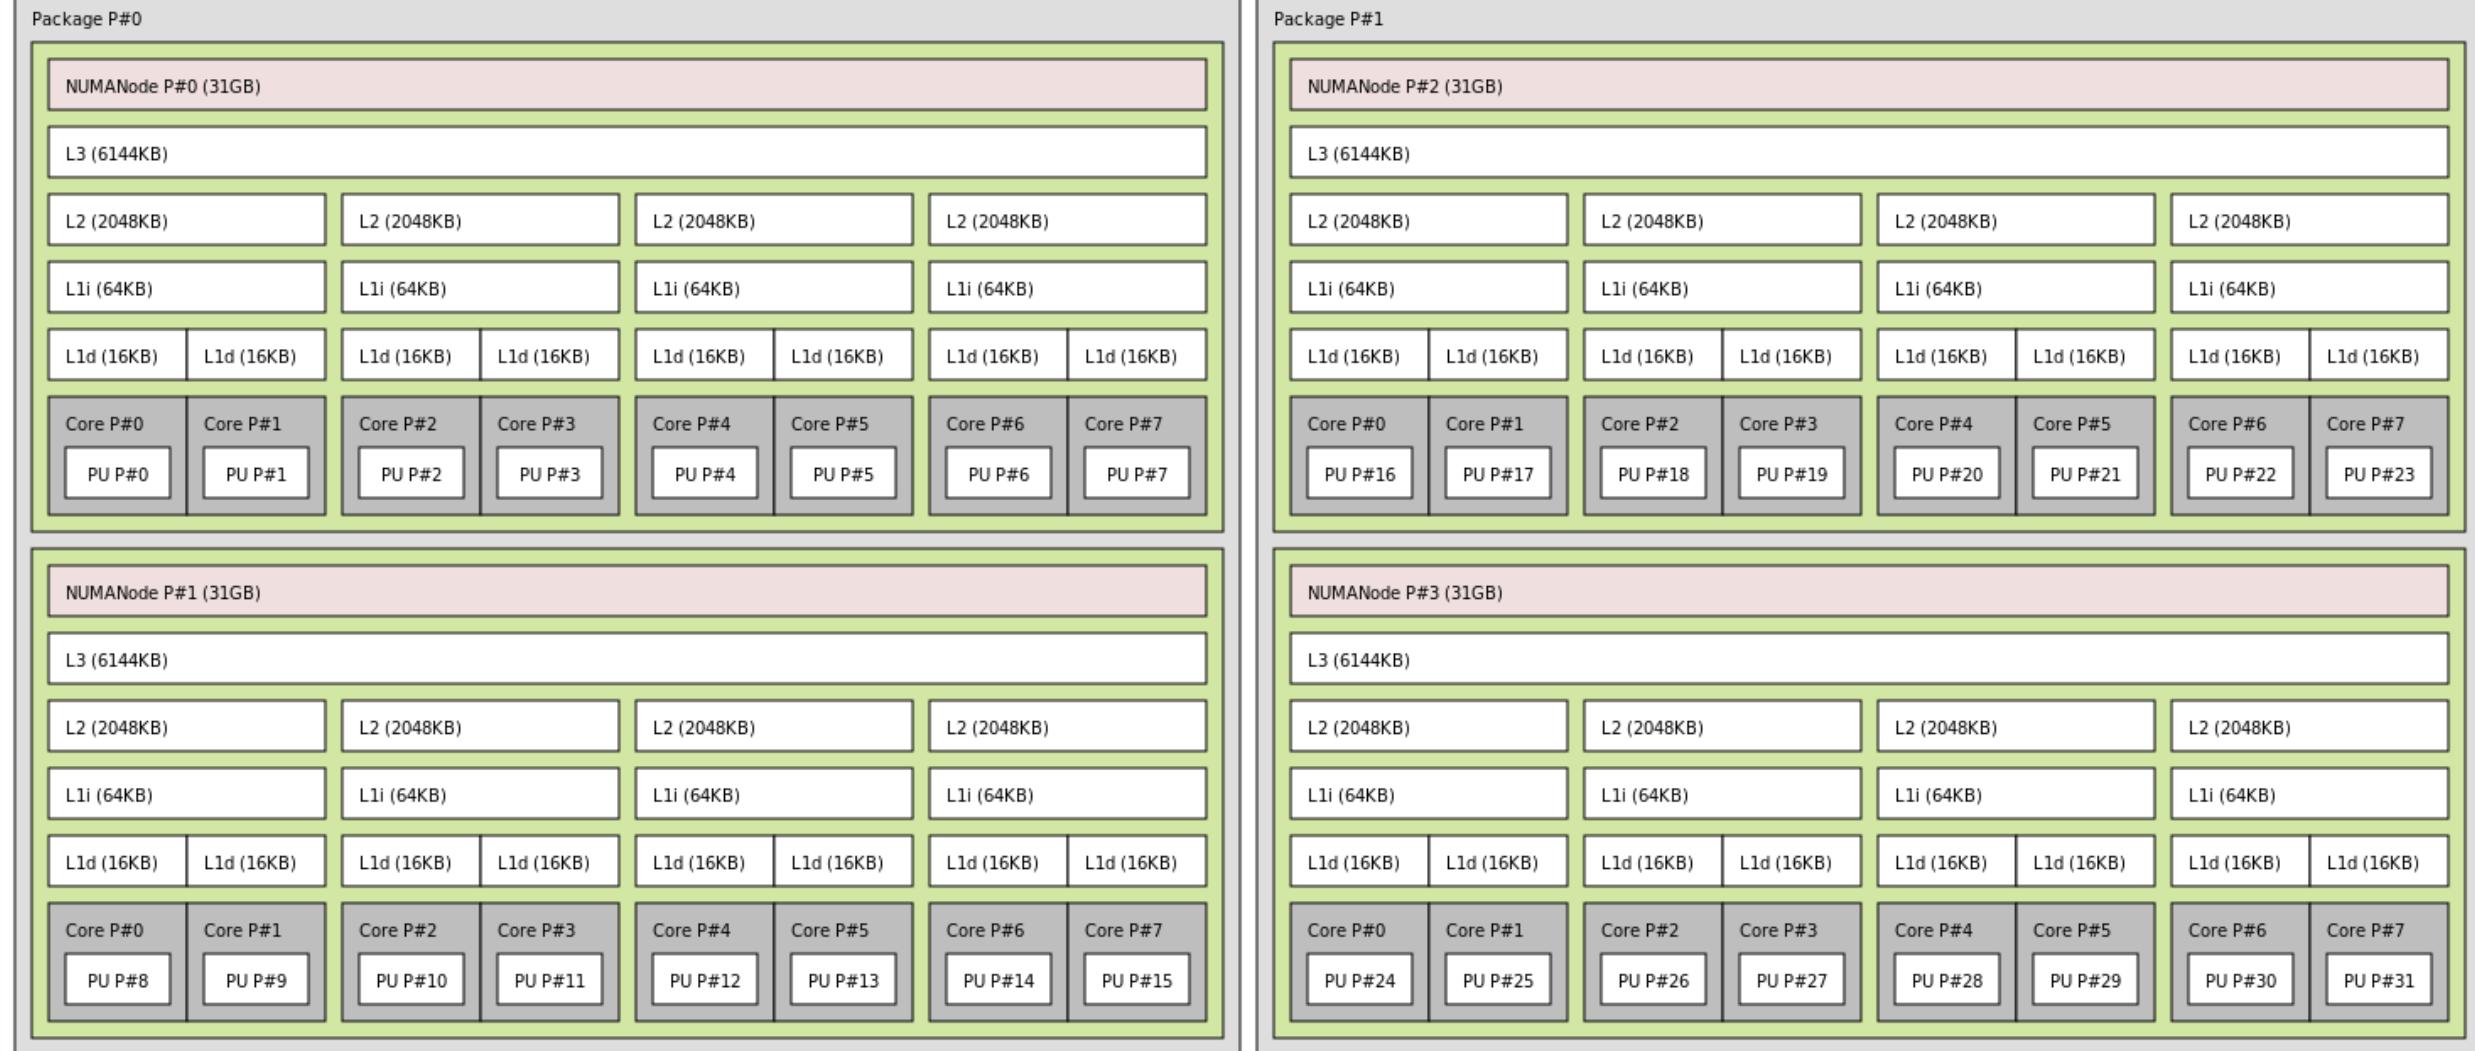
NVIDIA DGX A100 Station – AMD CPU



- **AMD 7742 ([Rome](#)) CPU, 64 cores in total**
- **hyper-treading on, 4 A100 GPUs attached**
- **8 dies & 16 core complexes**
- **creates additional layers in the NUMA hierarchy**

AMD 6386 (older) – clustered multi-threading

Machine (126GB total)



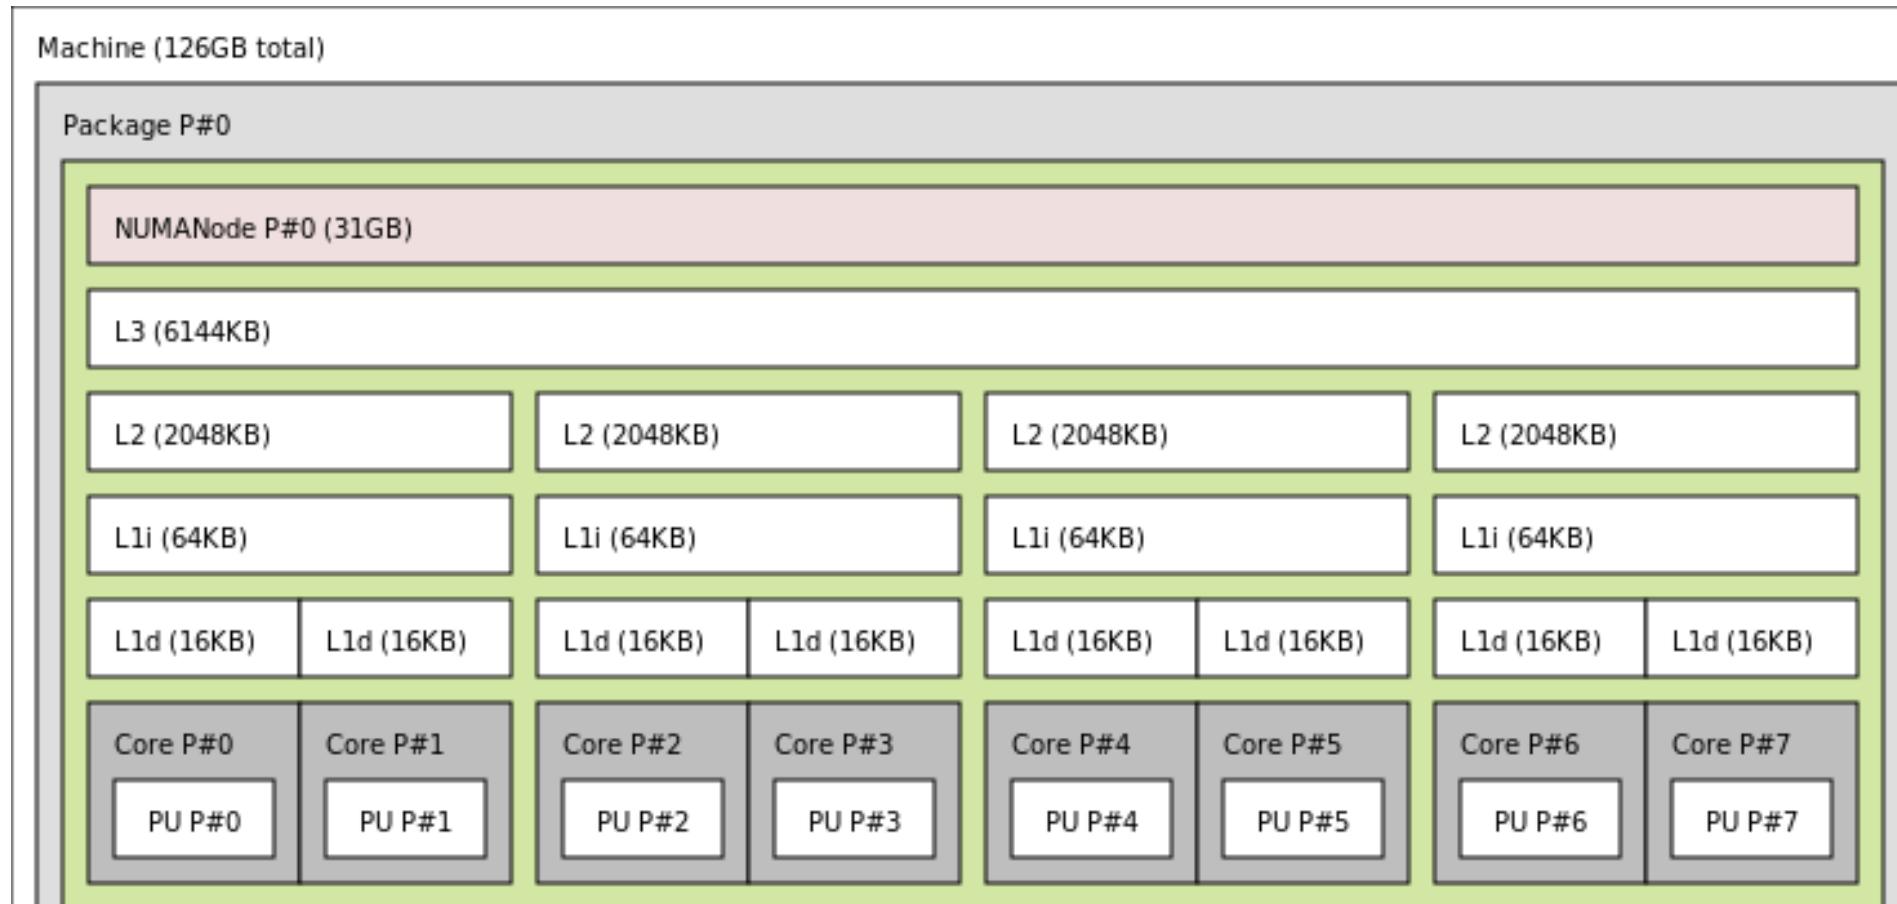
Host: dasyal

Indexes: physical

Date: Tue 04 Feb 2020 09:44:15 AM UTC

AMD 6386 (older) – clustered multi-threading

zoom in



agenda

lecture

- storage hierarchy overview
- local memory hierarchy: DRAM & caches
- hardware parallelism: implicit & explicit
- different memory hierarchies

exercises

- paper discussion
- c/cpp examples
- setting core affinity

experimental design case study

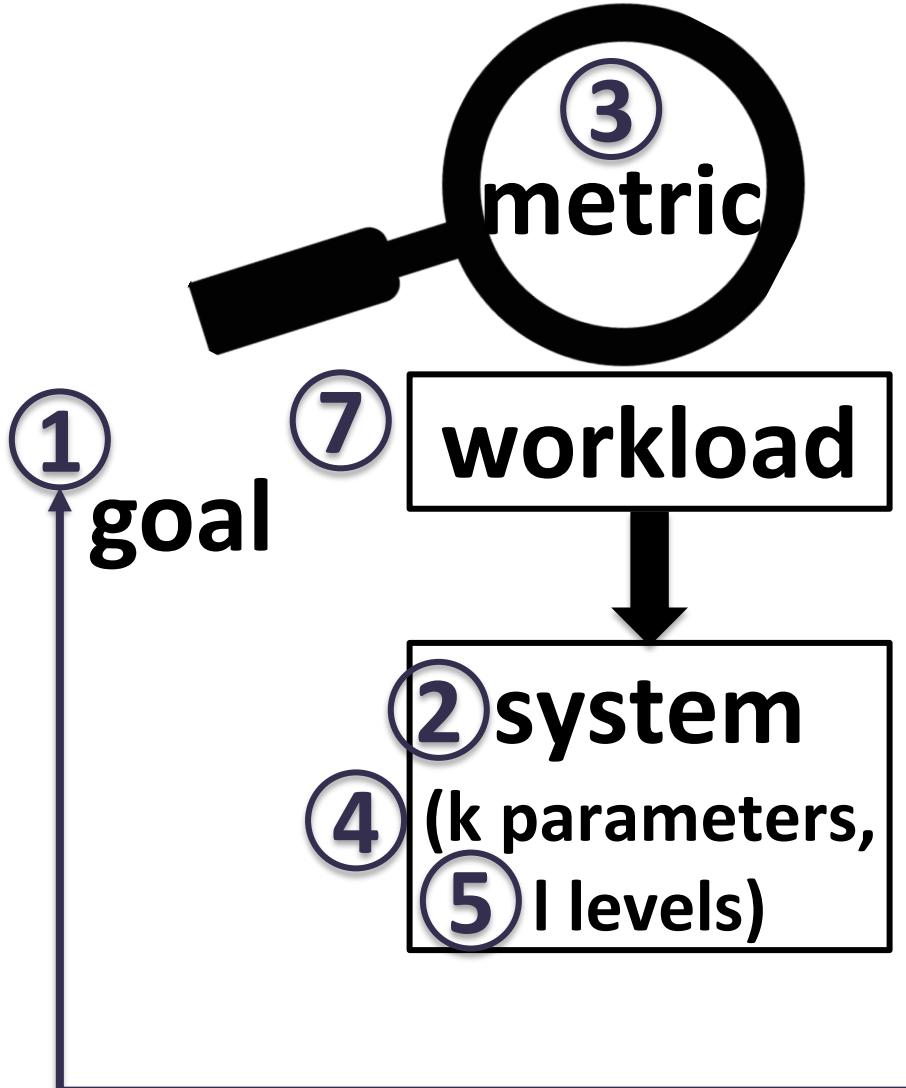
“Data Partitioning on Chip Multiprocessors”

analyze the experiments in this research paper
focusing on the following questions:

- what is the goal?
- what is/are the system(s)?
- what are the metrics?
- how many runs for the experiments?
- what are the good or bad points about the experiments?

also relevant for your first project!!

performance analysis – reminder



- experiment** 
- ⑥ decide on technique
(model, simulate, measure)
 - ⑧ actual experiment
 - ⑨ interpret results
 - ⑩ present results
- keep the big picture
in mind at all times!

goal: investigating how existing data partitioning mechanisms behave on multicore hardware. do old findings still hold?

system:

application-level: 4 partitioning mechanisms

OS-level: general purpose OS

hardware-level: multicore hardware

metrics: tuples per second (throughput), data TLB misses, data cache misses, metadata overhead

parameters (k): page size, hash bits (or #partitions), #threads,

levels(l): page size: 8KB, 64KB, 4MB, 256MB

hash bits: 1 ... 18, #threads: 1, 2, 4, 8, 16, 32

type of experiment: measurements (main), analytical approach to cache/tlb misses

workloads: with uniform data (8byte keys – 8byte payloads), 2^{24} tuples

experimental runs: 8 trials

useful commands

```
$ lscpu # cpu topology  
$ cat /proc/meminfo # checking hugepage size  
$ cat /proc/cpuinfo # core by core info  
$ cpuid # as above, but also TLB info
```